

Using State Abstractions to Compute Personalized Contrastive Explanations for AI Agent Behavior

Sarath Sreedharan, Siddharth Srivastava, Subbarao Kambhampati

CIDSE, Arizona State University, Tempe, AZ 85281 USA
{ ssreedh3, siddharths, rao } @ asu.edu

Abstract

There is a growing interest within the AI research community in developing autonomous systems capable of explaining their behavior to users. However, the problem of computing explanations for users of different levels of expertise has received little research attention. We propose an approach for addressing this problem by representing the user’s understanding of the task as an abstraction of the domain model that the planner uses. We present algorithms for generating minimal explanations in cases where this abstract human model is not known. We reduce the problem of generating an explanation to a search over the space of abstract models and show that while the complete problem is NP-hard, a greedy algorithm can provide good approximations of the optimal solution. We empirically show that our approach can efficiently compute explanations for a variety of problems and also perform user studies to test the utility of abstractions in explanations.

Keywords: Explanations for Plans, Abstractions, Contrastive Explanations

1. Introduction

AI systems have the potential to transform society by assisting humans in diverse situations ranging from extraplanetary exploration to assisted living. In order to achieve this potential, however, humans working with such systems need to be able to understand them just as they would understand human team members. This presents a number of challenges because most humans do not

7 understand AI algorithms and their behavior at the same intuitive level that
8 they understand other humans. Handling and possibly overcoming such knowl-
9 edge asymmetry requires us to develop and deploy AI systems that are capable
10 of providing cogent explanations for their actions/decisions to end users. A sig-
11 nificant challenge for any such system would be the fact that more often than
12 not, the AI system may be modeling and reasoning about the task with much
13 greater fidelity than the user is aware of (or capable of reasoning with). While
14 there have been a number of recent works on the problem of explaining plans
15 and actions chosen by agents (readers can refer to the survey [1] for previous
16 works in this direction), they have generally assumed that the user understands
17 the task at the same level of abstraction as the agent in question.

18 In this paper, we propose a new approach to this problem where the agent
19 explains its ongoing or planned behavior in a way that is both tailored to the
20 user’s background and is designed to reduce cognitive burden on the user’s end.
21 This is done by modeling a user’s expertise, or the level of detail at which a
22 user understands the task using abstracted models. We can estimate this level
23 based on questions that the user asks and provide explanations that are close
24 to this estimated level of expertise.

25 We consider explanations in the framework of counterfactual reasoning,
26 where a user who is confused by the agent’s activity (or proposed activity)
27 presents alternative behavior that they would have expected the agent to exe-
28 cute. This aligns with the widely held belief that humans expect explanations
29 to be *contrastive* [2]. In keeping with the terminology used in social sciences
30 literature, we will denote the set of alternative behaviors as *foils* to the proposed
31 robot behavior.

32 For instance, consider a mission-control operator who needs to manage an
33 autonomous robot on Mars in the midst of a sandstorm that could present
34 valuable data for analysis. If the robot proposes to go back to the base before
35 going to a vantage point for observing the storm, the operator would naturally be
36 perplexed, and may be motivated to ask the rover why it didn’t go directly to the
37 vantage point. While answering the operator’s queries it is important that the

38 explanation being given is tailored to meet the user’s background knowledge.
 39 Here an explanation that informs the operator of some specific mission goals
 40 that warranted this unintuitive plan, for example *“I am required to drop the*
 41 *collected samples in the base before going to the vantage point”*, may be preferred
 42 over a detailed explanation involving the specifics of the battery model or the
 43 rover motors (for example *“Motors of model #2310 needs to be recalibrated after*
 44 *every 20 miles and I need to go to base to recaliberate”*). As far as the rover
 45 is concerned, both of these explanations are equally valid reasons to choose the
 46 circuitous route, but a mission control operator may find the former easier to
 47 understand while an engineer may better appreciate the latter. This level of user
 48 specificity requires methods that estimate possible models that can capture the
 49 user’s level of understanding of the task. As mentioned, we will make use of
 50 the questions (i.e the foils) raised by the user for the specific task at hand (and
 51 potentially even the history of previous interactions) to build such estimates.
 52 Accurate estimate of the user’s expertise not only lets us control the level of
 53 detail in explanation but also allows us to provide the most concise explanation
 54 (by avoiding unnecessary details) and thereby reduce the cognitive burden of
 55 the user.

56 In this paper we present the **Hierarchical Expertise-Level Modeling**
 57 or the **HELM** approach for facilitating such context and user-specific explana-
 58 tions. We assume that the user’s understanding of the task is an abstraction
 59 of the model used by the robot; which captures both the limited information
 60 and computational capabilities of the user. HELM generates appropriate ex-
 61 planations by searching through a *model lattice* of possible abstractions of the
 62 agent’s model. The model lattice provides a concise way for the system designer
 63 to encode their prior knowledge about potential users. Each model within this
 64 lattice represents a different level of understanding of the task, with the high-
 65 est fidelity representation (corresponding to the most detailed understanding of
 66 the domain used by the robot) forming the base of the lattice and the model
 67 representing the most naive understanding of the task (for example one held by
 68 a lay person) forming the highest nodes. Since the user’s level of expertise is

69 unknown to the agent, it has to estimate the human model before searching for
70 an explanation.

71 We focus on contrastive explanations, where an explanation that is an answer
72 to a question of the form “Why P and not Q?”, in our case, P and Q are stand-
73 ins for the current robot plan and the foil respectively. Most existing works in
74 explanation for plans have focused on answering the first part “Why P?” (for
75 example works like [3, 4] have looked at identifying causal explanations for each
76 action), so the majority of this work will focus on finding a concise explanation
77 for the latter part of the full question. Specifically, our explanations will consist
78 of model information that may be absent in the user’s abstract model and
79 possible proofs for foil failure. Thus, in addition to helping convince the user
80 of the incorrectness of the foils in question, the explanations should also shift
81 the user’s model to a more accurate model in the lattice. This approach could
82 be understood as a variant of the model refinement methods discussed in the
83 counter-example guided model checking (CEGAR) literature [5]. Our methods
84 extend these principles to settings with uncertainty regarding the current level
85 of abstraction of the model (a non-issue in the model-checking settings where
86 CEGAR methods are typically used).

87 This paper generalizes and extends our recent work [6] with extended theo-
88 retical and empirical results and exposition. In addition to clarifying the con-
89 cepts our contributions include

- 90 • We consider the use of non-standard lattices as a way to allow designers
91 to incorporate more information about the user’s model in to the expla-
92 nation generation process and discuss potential computational tradeoff
93 introduced by the use of such lattice types over the ones considered in the
94 rest of the paper.
- 95 • We investigate the use of such methods for domains that contain state de-
96 pendent costs (hence affected by the abstraction) and discuss the potential
97 explanatory dialogue that could occur in such settings.
- 98 • We also show how our method could be used in cases where the user

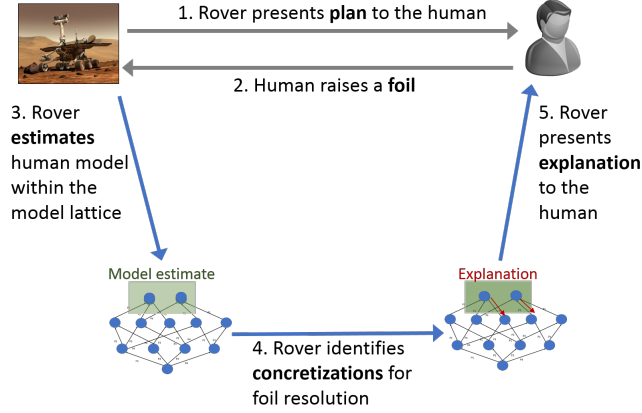


Figure 1: An illustration of the hierarchical explanation process. The human observer who views the task at a higher level of abstraction expects the rover to execute a different plan from the one chosen by the rover. The rover presents the human with an explanation it believes will help resolve the foils in the human’s updated model.

99 model may not just be abstract but the user may also hold erroneous
 100 beliefs about the task.

- 101 • We perform a user study to verify the utility of abstraction in generating
 102 explanation that are easier for users to work with.

103 The rest of this paper is structured as follows. Section 2 a brief overview
 104 of the background and in Section 3 we present our formal framework. Section
 105 4 covers different approaches for generating explanations and sections 4.1, 4.2
 106 and 4.3 extend these methods to more general settings. Section 5 presents
 107 evaluations of the method. In sections 6 and 7, we will discuss the related work
 108 and possible future directions.

109 2. Background

110 In this work, we focus on abstractions that form models by projecting out
 111 state fluents. While the presentation in the following sections is equally valid
 112 for both predicate and propositional abstractions, we will focus on propositional

113 abstractions to keep our formulation clear and concise and later discuss potential
 114 changes required to meet the requirements of predicate abstractions. We will
 115 look at planning models of the form $\mathcal{M} = \langle P, S, A, I, G \rangle$ where P gives the set
 116 of state fluents, S the set of possible states, A the set of actions, I the initial
 117 state and G the goal. Each state $s \in S$ is uniquely represented by the set of
 118 propositions that are true in that state, i.e, $s \subseteq P$.

119 Each action $a \in A$ is associated with a set of positive preconditions prec_a^+
 120 (specified as a conjunction of propositions) and negative preconditions prec_a^-
 121 that need to hold for the effects (e_a) of that action to be applied to a particular
 122 state. Each effect set e_a can be further separated into a set of add effects e_a^+
 123 and a set of delete effects e_a^- . The result of executing an action a on a state s
 124 in this setting is defined as $a(s) = (s \cup e_a^+) \setminus e_a^-$, if $\text{prec}_a^- \subseteq s \wedge \text{prec}_a^+ \cap s = \emptyset$.
 125 A plan π is defined as a sequence of actions $\langle a_1, \dots, a_n \rangle$, n being the size of the
 126 plan), and a plan is said to solve \mathcal{M} (i.e, $\pi(I) \models_{\mathcal{M}} G$) if $\pi(I) \supseteq G$.

127 Automated planning has a long tradition of employing abstraction both for
 128 plan generation (cf. [7]) and for generating heuristics (cf. [8, 9]) and a number
 129 of different abstraction schemes have been proposed in these works. In fact,
 130 state abstractions as presented in this work have been widely used in pattern
 131 databases and are referred to as projections in that literature (cf. [10, 11]).
 132 Following works like [8, 12], we will also use the concept of a transition sys-
 133 tem induced by the planning model to define state abstractions. Intuitively, a
 134 transition system constitutes a graph where the nodes represent possible states,
 135 and the edges capture the transitions between the states that are valid in the
 136 corresponding planning model.

137 Formally a transition system \mathcal{T} corresponding to a model \mathcal{M} can be repre-
 138 sented by a tuple of the form $\mathcal{T} = \langle S, L, T, s_0, S_g \rangle$, where S is the set of possible
 139 states in \mathcal{M} , L is the set of transition labels (corresponding to the action that
 140 induce that transition), T is the set of possible labeled transitions, s_0 is the
 141 initial state and S_g is the set of states that satisfies the goal specified by \mathcal{M} .
 142 We will refer to \mathcal{T} to be the *safe transition* system induced by a model \mathcal{M} , if
 143 and only if, for any labeled transition $\langle s, a, s' \rangle \in T$, we have $\text{prec}_a^+ \subseteq s$ and

144 $\text{prec}_a^- \cap s = \emptyset$. Through most of this work we will focus our attention on cases
 145 where the semantics of the planning task is defined in terms of safe transition
 146 systems.

147 **Definition 1.** A propositional abstraction function f_Λ for a set of propositions
 148 Λ and state space S , defines a surjective mapping of the form $f_\Lambda : S \rightarrow X$,
 149 where X is a projection of S , such that for every state $s \in S$, there exists a
 150 state $f_\Lambda(s) \in X$ where $f_\Lambda(s) = s \setminus \Lambda$.

151 **Definition 2.** For a planning model $\mathcal{M} = \langle P, S, A, I, G \rangle$ with a corresponding
 152 transition system \mathcal{T} , a model $\mathcal{M}' = \langle P', S', A', I', G' \rangle$ with a transition system
 153 \mathcal{T}' is considered an **abstraction of \mathcal{M}** for a set of propositions Λ , if for every
 154 transition $s_1 \xrightarrow{a} s_2$ in \mathcal{T} corresponding to an action a , there exists an equiv-
 155 alent transition $f_\Lambda(s_1) \xrightarrow{a'} f_\Lambda(s_2)$ in \mathcal{T}' , where a' is part of the new action set A' .
 156

157 We will slightly abuse notation and extend the abstraction functions to mod-
 158 els and actions, i.e in the above case, we will have $\mathcal{M}' \in f_\Lambda(\mathcal{M})$ (where $f_\Lambda(\mathcal{M})$
 159 is the set of all models that satisfy the above definition for the set of fluents Λ)
 160 and similarly we will have $a' \in f_\Lambda(a)$. As per Definition 2, the abstract model
 161 is *complete* in the sense that all plans that were valid in the original model will
 162 have an equivalent plan in this new model. We will use the operator \sqsubset to cap-
 163 ture the fact that the model \mathcal{M}' is an abstraction of \mathcal{M} , i.e if $\mathcal{M} \sqsubset \mathcal{M}'$ then
 164 there exist a set of propositions Λ such that $\mathcal{M}' \in f_\Lambda(\mathcal{M})$.

165 2.1. Designing Complete Abstractions

166 While there exists a number of works that have looked at the problem of
 167 designing abstractions (cf. [13, 7, 12]), unfortunately many of these works have
 168 considered directly updating transition system or using specialized or more ex-
 169 pressive problem formulation to capture abstract models. Thankfully, the fact
 170 that we are interested in complete abstractions (as opposed to sound abstrac-
 171 tions) means we can employ simpler model transformation schemes to generate

abstract models. In particular, we will consider transformations that simply drops the set of literals to be abstracted from all the action definitions, i.e,

Theorem 1. *For a given model $\mathcal{M} = \langle P, S, A, I, G \rangle$ and a set of propositions Λ , a model $\mathcal{M}' = \langle P', S', A', I', G' \rangle$ is a complete abstraction under safe execution semantics for Λ , if $P' = P - \Lambda$, $S' = [S]_{f_\Lambda}$, $I' = f_\Lambda(I)$, $G' = G \setminus \Lambda$ and for every $a \in A$ (where $a = \langle prec_a^+, prec_a^-, eff_a^+, eff_a^- \rangle$) there exists $a' \in A'$, such that $a' = \langle prec_a^+ \setminus \Lambda, prec_a^- \setminus \Lambda, eff_a^+ \setminus \Lambda, eff_a^- \setminus \Lambda \rangle$.*

Proof Sketch. To see why the new model would be an complete abstraction, consider a transition $\langle s, a, s' \rangle$ induced by \mathcal{M} . Now as per the definitions of safe transition systems, we know that $s \subseteq prec_a^+$ and $s \cap prec_a^- = \emptyset$ and $s' = s \setminus eff_a^- \cup eff_a^+$. Its easy to see that given this setting, $(s \setminus \Lambda) \subseteq (prec_a^+ \setminus \Lambda)$ and $(s \setminus \Lambda) \cap (prec_a^- \setminus \Lambda) = \emptyset$, which means there must be an action $a' \in A'$ that is executable in $f_\Lambda(s)$. Similarly we can show the result of executing a' must be $f_\Lambda(s)$, this shows that \mathcal{M}' is a complete abstraction of \mathcal{M} as every transition induced by it is present in the transition system induced by \mathcal{M}' . \square

An important point to note here is that this transformation scheme generates a unique abstract model for each model and proposition set, and we will denote this unique model as $f_\Lambda(\mathcal{M})$. For the rest of the paper, we will mainly focus on this method to induce the abstractions, but general framework of explanation generation discussed in this paper can be adapted to other methods of generating abstract models. In cases, where we prove specific results or present optimization that rely on this abstraction procedure we will denote the abstraction function by f_Λ^{safe} to differentiate it from other methods. With the definition of abstraction and related notations in place, we will look at our explanatory setting and a way to capture the space of possible user models that would allow for efficient estimation of unknown user model given user queries. While the above operation is defined for propositional fluents, we can perform similar operations on the lifted domain, where projecting out a predicate would correspond to projecting out a set of propositional fluents from the grounded domain.

3. Hierarchical Expertise-Level Modeling

As mentioned earlier, we are investigating explanatory settings where the user’s understanding of the task can be represented as an abstraction of the robot’s model. While the exact level of abstraction may be unknown, given a set of candidate state fluents that may be missing from the human model, we can capture the potential models and their relationship through a **model lattice**

Definition 3. For a model $\mathcal{M}^\#$, the *model lattice* \mathcal{L} is a tuple of the form $\mathcal{L} = \langle \mathbb{M}, \mathbb{E}, \mathbb{P}, \ell \rangle$, where \mathbb{M} is the set of lattice nodes such that $\mathcal{M}^\# \in \mathbb{M}$ and $\forall \mathcal{M}' \in \mathbb{M}, \mathcal{M}^\# \subseteq \mathcal{M}'$, \mathbb{E} is the lattice edges, \mathbb{P} is the superset of propositions considered for abstraction within this lattice and ℓ is a function mapping edges to labels. Additionally, for each edge $e_i = (\mathcal{M}_i, \mathcal{M}_j)$ there exists a proposition $p \in \mathbb{P}$ such that $f_p(\mathcal{M}_i) = \mathcal{M}_j$ and $\ell(\mathcal{M}_i, \mathcal{M}_j) = p$.

Thus each edge in this lattice corresponds to an abstraction formed by projecting out a single proposition (represented by the label of the edge). We can also define a concretization function γ_p that retrieves the model that was used to generate the given abstract model by projecting out the proposition p , i.e., $\gamma_p(\mathcal{M}) = \mathcal{M}'$ if $(\mathcal{M}', \mathcal{M}) \in \mathbb{E}$ and $\ell(\mathcal{M}', \mathcal{M}) = p$ else $\gamma_p(\mathcal{M}) = \mathcal{M}$.

For a given lattice, if each node in \mathbb{M} has an incoming edge for every proposition missing from its corresponding model then we will refer to such lattices as being *Proposition Conserving* lattices.

Definition 4. A lattice \mathcal{L} is *proposition conserving*-iff for any model $\mathcal{M} \in \mathbb{M}$ ($\mathcal{M} = \langle P_{\mathcal{M}}, S_{\mathcal{M}}, A_{\mathcal{M}}, I_{\mathcal{M}}, G_{\mathcal{M}} \rangle$) and $\forall p \in \mathbb{P}$, if p is not in $P_{\mathcal{M}}$ then there exists a model $\mathcal{M}' \in \mathbb{M}$, such that $(\mathcal{M}', \mathcal{M}) \in \mathbb{E}$ and $\ell(\mathcal{M}', \mathcal{M}) = p$.

Notice that enforcing conservation of propositions doesn’t require any further assumptions about the human model and can be easily ensured while generating the lattice. Additionally, we will call a proposition conserving lattice that contains an abstract node corresponding to each possible subset of \mathbb{P} as the

230 *Complete Abstraction Lattice* for \mathcal{M} given \mathbb{P} . The earlier parts of this paper
 231 will assume a proposition conserving lattices as they will allow us to simplify
 232 discussions and provide efficient solutions. In later sections, we will relax these
 233 assumptions and will look at potential tradeoffs for using non-proposition con-
 234 serving lattices.

235 We also assume that all abstraction functions used in generating the models
 236 in the lattice are commutative and idempotent, i.e., $f_{p_2}(f_{p_1}(\mathcal{M})) = f_{p_1}(f_{p_2}(\mathcal{M}))$
 237 and $f_{p_1}(f_{p_1}(\mathcal{M})) = f_{p_1}(\mathcal{M})$. In the wider literature, a lattice is generally defined
 238 to have a unique maximal element and a unique minimal element. While the
 239 abstraction lattices we consider in this work will have a unique minimal element
 240 (i.e the most concrete nodes), we do not assume that the lattices have a single
 241 maximal node (Figure 2 presents an example lattice that does not have a unique
 242 maximal node), in that sense the abstraction lattice may be better understood
 243 as meet-semilattices, but we will use the term model lattice or abstraction lattice
 244 for convenience.

245 As mentioned earlier, we consider an explanation generation setting where
 246 the human observer (H) uses a task model (this model will be denoted as $\mathcal{M}_H =$
 247 $\langle P_H, S_H, A_H, I_H, G_H \rangle$), that is a more abstract version of the robot’s model
 248 ($\mathcal{M}_R = \langle P_R, S_R, A_R, I_R, G_R \rangle$). While the robot may not know \mathcal{M}_H , it knows
 249 that \mathcal{M}_H is a member of the set \mathbb{M} for the lattice \mathcal{L} . The human comes up
 250 with a **foil set** $\mathbf{F} = \{\pi_1, \pi_2, \dots, \pi_m\}$ that the robot needs to refute by providing
 251 an explanation regarding the task. The explanation should contain information
 252 about specific domain properties (i.e., state fluents) that are missing from the
 253 human’s model, how these properties affect different actions (For example, which
 254 actions use these propositions as preconditions and which ones generate/delete
 255 them) and how the inclusion of these fluents result in the invalidity of the given
 256 foils. To illustrate the utility of such explanations consider an example involving
 257 a simplified version of the rover domain mentioned earlier.

258 **Example 1.** *Let us suppose that the rover uses a modified version of the IPC*
 259 *rover domain [14] that also takes into account the battery level of the rover. Each*

260 rover operation has a different energy requirement, and the battery level needs to
 261 be above a predefined threshold for it to execute them, e.g., it can perform rock
 262 sampling only if the battery level is above 75%. Furthermore, the rover needs to
 263 visit the base station (i.e., the lander) and perform a reset action to recharge its
 264 batteries.

265 The rover knows that the human observer is at most ignorant of its energy
 266 requirements, ability to use solar cells and/or storage capabilities. So the model
 267 lattice \mathcal{L} needs to consider abstractions corresponding to the following propo-
 268 sitions

269 $\mathbb{P} = \{ \text{battery_level_above_25_perc}, \text{battery_level_above_50_perc},$
 270 $\text{battery_level_above_75_perc}, \text{full_store1}, \text{solar_panels_activated} \}.$

Figure 2 shows the lattice that the robot would use in this setting. Here we
 will create each abstract model by following the process discussed in section 2.1.
 For example, consider the action `sample_rock_store0.w1`, it has the following defi-
 nition

$$\langle \{ \text{battery_level_above_75_perc}, \text{at_w1}, \text{empty_store1}, \text{has_store_store1} \}, \{ \}, \\ \{ \text{full_store1}, \text{has_rock_sample} \}, \{ \text{empty_store1}, \text{battery_level_above_75_perc} \} \rangle$$

271

Now in an abstract version of this model, if the propositions `full_store1`, `battery_level_above_75_perc` are dropped the definition becomes

$$\langle \{ \text{at_w1}, \text{has_store_store1} \}, \{ \}, \\ \{ \text{has_rock_sample} \}, \{ \text{empty_store1} \} \rangle$$

272

Here the robot presents the plan

$$\pi_R = \langle \text{navigate_w0_lander}, \text{reset_at_lander}, \\ \text{navigate_lander_w1}, \text{sample_rock_store0_w1} \rangle$$

and a naive observer may respond by proposing the foil set with a single plan

$$F = \langle \text{navigate_w0_w1}, \text{sample_rock_store0_w1} \rangle$$

If the observer was an engineer, they might instead raise a foil that already takes into account the energy requirements

$$F' = \{ \langle \text{navigate_w0_w1}, \\ \text{receive_energy_from_solar_cells}, \text{sample_rock_store0_w1} \rangle \}$$

273 If the robot knew that the human was ignorant about all the battery level
 274 predicates and nothing else, the robot could help resolve the naive foil by in-
 275 forming them about the fact that action sample rock requires the battery to be
 276 above 75% (i.e describing the proposition `battery_level_above_75_perc`). In terms
 277 of the human model, this would involve setting the value of the proposition `battery_level_above_75_perc` false in the initial state, updating the precondition of
 278 `sample_rock_store0_w1` to include the fact (among other actions) and adding it
 279 as an add effect to the action `reset_at_lander`. In this updated model the human
 280 foil can no longer achieve the goal. In the case, of expert foil, the robot would
 281 need to inform the user about the proposition `solar_panels_activated` and that
 282 the action `receive_energy_from_solar_cells` require the solar panels to be activated
 283 which is not true for the rover. Thus in each case explanations to be provided to
 284 user can be generated once we know the set of propositions whose concretization
 285 is required to refute the given foils (henceforth referred to as *explanatory fluent*
 286 *set*).

Definition 5. Let $E = \{p_1, \dots, p_n\}$ be a set of fluents, then E is said to be an explanatory set for the human model \mathcal{M}_H and a foil set F if

$$\forall \pi \in F, \pi(I_{\gamma_E(\mathcal{M}_H)}) \not\models_{\gamma_E(\mathcal{M}_H)} G_{\gamma_E(\mathcal{M}_H)}$$

288 Where $\gamma_E(\mathcal{M}_H)$ is the model obtained by applying the concretizations corre-
 289 sponding to E on the model \mathcal{M}_H .

290 In the case of projection based abstractions of the form defined in Section
 291 2.1, we can directly provide the model components covered by the explanatory
 292 fluent set as part of the final explanatory message provided to the user. For
 293 other abstraction techniques, we may need to employ more specialized methods

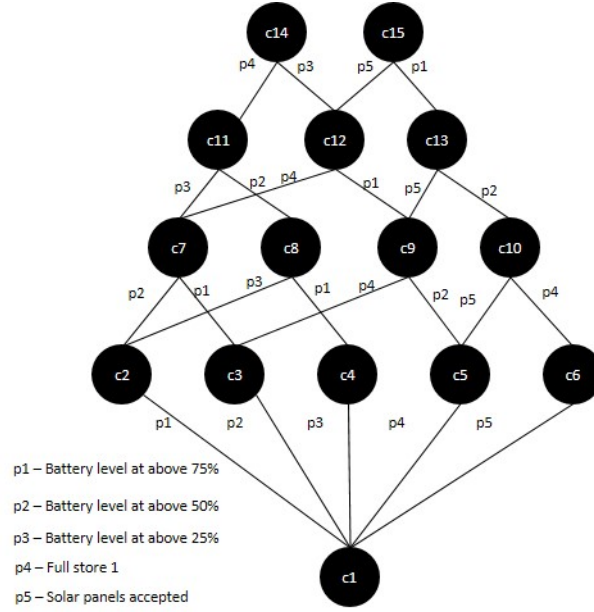


Figure 2: A possible abstraction lattice for the rover domain.

294 to generate explanatory messages from the fluents. In Example 1 if we are to
 295 focus on the naive foil, the rover would have difficulty coming up with a single
 296 explanation as it does not know \mathcal{M}_H . However, it can restrict its attention to
 297 just the models that are consistent with the foils. In this scenario, it would
 298 correspond to $\{c2, c7, c8, c11, c12, c14, c15\}$.

299 Now we need to find a way of generating sets of explanatory fluents given
 300 this reduced set of models.

301 **Proposition 1.** *Let \mathcal{M}_i be some model in \mathcal{L} such that $\mathcal{M}_H \sqsubseteq \mathcal{M}_i$. If E is*
 302 *explanatory for \mathcal{M}_i and some foil set F , then E must also explain F for \mathcal{M}_H .*

303 This proposition directly follows from the fact that for a proposition con-
 304 serving lattice $\gamma_E(\mathcal{M}_i)$ will be a logical weaker model than $\gamma_E(\mathcal{M}_H)$. Next, we
 305 will define the concept of a minimal abstracting set for a given lattice \mathcal{L} and
 306 foils F

307 **Definition 6.** Given an the abstraction lattice $\mathcal{L} = \langle \mathbb{M}, \mathbb{E}, \mathbb{P}, \ell \rangle$ and a foil set

308 F , the *minimal abstracting set* \mathbb{M}_{min}^F is the maximal elements of the subset of
 309 all the models that are consistent with F .
 310 $\mathbb{M}_{min}^F = \{\mathcal{M}_i | \mathcal{M}_i \text{ is a maximal element of } \mathbb{M}_{sat}\}$ where $\mathbb{M}_{sat} = \{\mathcal{M} \mid \mathcal{M} \in$
 311 $\mathbb{M}, \forall \pi \in F(\pi(I_{\mathcal{M}}) \models_{\mathcal{M}} G_{\mathcal{M}})\}$

312 **Proposition 2.** *For a given model lattice \mathcal{L} , the minimal abstracting set \mathbb{M}_{min}^F*
 313 *is a subset of the maximal elements of the entire abstraction lattice.*

314 The above property ensures that when searching for the minimal abstracting
 315 set, we do not need to test the entire set of nodes or even need to know the
 316 entire lattice. In Example 1, the minimal abstracting set for the naive foil will
 317 be $\mathbb{M}_{min}^F = \{c14, c15\}$.

318 If we can find an explanation that is valid for all the models in \mathbb{M}_{min}^F then
 319 by Proposition 1 it must work for \mathbb{M}_H as well.

Proposition 3. *For a given model lattice \mathcal{L} and a set of foils F and the min-
 imal abstraction set \mathbb{M}_{min}^F , there exists an explanatory fluent set E such that
 $\forall \mathcal{M}' \in \mathbb{M}_{min}^F$ and $\forall \pi \in F$,*

$$\pi(I_{\gamma_E(\mathcal{M}')}) \not\models_{\gamma_E(\mathcal{M}')} G_{\gamma_E(\mathcal{M}')}$$

320 It is easy to see why this property holds, as any explanation that involves
 321 concretizing all possible propositions in \mathbb{P} satisfies this property.

322 In most cases, we would prefer to compute the minimal or cheapest expla-
 323 nation to communicate. If all concretizations are equally expensive to commu-
 324 nicate to the explainee, then this would correspond to finding the explanatory
 325 fluents set with the smallest size. For the naive foil in the rover example, even if
 326 the human is unaware of multiple task details, the robot can easily resolve the
 327 explainee's doubts by just explaining the concretizations related to the proposi-
 328 tion `battery_level.above.75_perc` without getting into other details. Describing the
 329 details of remaining propositions is unnecessary and in the worst case might
 330 leave the human feeling overwhelmed and confused. In this case, the explana-
 331 tion would just include information regarding battery levels and how to identify

332 when the battery level is or above 75% and model updates like

333 sample_rock-has-precondition-battery_level_above_75_perc

334 sample_soil-has-precondition-battery_level_above_75_perc

335 ...

336 Before delving into the optimization version of the problem, let us look at the
337 complexity of the corresponding decision problem

338 **Theorem 2.** *Given a the set of foils F and the corresponding minimal abstrac-*
339 *tion set \mathbb{M}_{min}^F for a model \mathcal{M} , the problem of identifying whether an explanatory*
340 *fluent set of size k exists for the complete lattice (which is not given) defined*
341 *over an abstraction function f is **NP-complete**, provided the abstract func-*
342 *tion generates planning models that belong to the class described in Section 2 in*
343 *polynomial time.*

344 *Proof (Sketch).* The fact that we can test the validity of the given explanation
345 in polynomial time (size of the explanation is guaranteed to be smaller than $|\mathbb{P}|$)
346 shows that the problem is in **NP**. We can show **NP-completeness** by reducing
347 the set covering problem [15] to an instance of the explanation generation prob-
348 lem. Let's consider a set covering problem with U as the universe set and S as
349 the set of sub-collections. Now let us create an explanation generation problem
350 where the set of foils F is equal to U and the propositions in the set \mathbb{P} contain
351 a proposition for each member of S . Additionally concretizing with respect to
352 a proposition will resolve only the foils covered by its corresponding subset in
353 S . For this setting, the \mathbb{M}_{min}^F consists of a single node that contains none of the
354 propositions (and hence all the foils hold) and the concrete model contains all
355 of them. Now if we can come up with a set of explanatory fluents of size k in
356 this setting, then this explanation corresponds to a set cover of size k . \square

357 The above result considers a case where the lattice needs to be generated
358 on the fly from the minimal abstraction set. Though there may be cases where
359 the designer may be able to provide an explicit and smaller non-proposition
360 conserving lattice upfront. As we will see in Section 4.1, such lattices can be
361 used to capture the designer's knowledge about the end-users.

362 4. Generating Optimal Explanations

363 As mentioned earlier, we are interested in producing the minimal explana-
 364 tion. Additionally, in most domains, the cost of communicating the concretiza-
 365 tion details could vary among propositions. An explanation that involves a
 366 proposition that appears in every action definition might be harder to commu-
 367 nicate than one that only uses a proposition that is part of the definition of a
 368 single action.

369 In addition to the actual size, the comprehensibility of the explanations
 370 may also depend on factors like human’s mental load, the familiarity with the
 371 concepts captured by the propositions, etc.. To keep our discussions simple,
 372 we will restrict the cost of communicating an explanation to just the number
 373 of unique model updates this explanation would bring about in the human
 374 model. We will use the symbol $C_p^{\mathcal{E}}$ to represent the cost of communicating the
 375 changes related to the proposition p and also overload it to be applicable over
 376 sets of propositions.

377 Now our problem is to find the optimal explanation (represented as E_{min})
 378 for a given set of foils F or more formally

379 **Definition 7.** *A set of fluents E is said to be the optimal explanatory fluent*
 380 *set for the human model \mathcal{M}_H and a foil set F , if*

- 381 1. *if E is an explanatory set and*
- 382 2. *there exists no other set \hat{E} , such that \hat{E} is also an explanatory set and*
 383 $C_E^{\mathcal{E}} > C_{\hat{E}}^{\mathcal{E}}.$

384 Given the fact that the human model is not known to start with, it may
 385 appear that there is no way to generate optimal explanations for the human
 386 model directly. A possible alternative might be to try identifying the set of
 387 fluents that is optimal for the set of models that could be \mathcal{M}_H . Calculating such
 388 an explanation naively could be extremely expensive as identifying all possible
 389 candidates for the human model would involve testing each node in the lattice
 390 for whether its a potential candidate for the human model and then searching

391 over the space of all explanatory fluent set to find one that is optimal for the
 392 entire set of candidate models (where the optimality for a set of models is defined
 393 to be the cheapest set of fluents that is explanatory for all the models in the set).
 394 Thankfully, the properties of the lattice allow us to compute optimal solutions
 395 without keeping track of the entire set. Moreover, for lattices containing abstract
 396 models generated using procedures discussed in Section 2.1, we will see how
 397 fluent sets that are optimal for minimal abstracting set are still optimal for
 398 the original human model. *That is uncertainty over human models results in no*
 399 *loss of optimality.* But before proving that property, we will define the idea of
 400 the *resolution set*, that captures the specific plans resolved by concretizing the
 401 given propositions (i.e the proposition appears as an unsatisfied precondition or
 402 goal in the plan).

Definition 8. For a set of models \mathbb{M}' , a foil set F and a proposition p , the
resolution set $\mathcal{R}_F(\mathbb{M}', p)$ gives the subset of foils that no longer holds in the
 concretized models generated through $f_{\Lambda}^{\text{safe}}$, i.e

$$\begin{aligned}
 \mathcal{R}_F(\mathbb{M}', p) = \{ \pi \mid \pi \in F \wedge (\forall \mathcal{M}' \in \mathbb{M}' (\pi(I_{\gamma_p(\mathcal{M}')})) \not\models_{\gamma_p(\mathcal{M}')} G_{\gamma_p(\mathcal{M}')} \wedge \\
 \pi(I_{\mathcal{M}'}) \models_{\mathcal{M}'} G_{\mathcal{M}'})) \}
 \end{aligned}$$

403 The idea of generating resolution sets are again closely related to the idea of
 404 resolving counter-examples used in CEGAR based method. We will also use \mathcal{R}_F
 405 to also represent the set of foils resolved by a set of propositions. For notational
 406 convenience, we will use $\mathcal{R}_F(\mathbb{M}', \{\})$ to capture the subset of foils that do not
 407 hold in the current model set \mathbb{M}' .

Proposition 4. For a set of model \mathbb{M}' and a foil set F

$$\mathcal{R}_F(\mathbb{M}', \{p_1, p_2\}) = \mathcal{R}_F(\mathbb{M}', \{p_1\}) \cup \mathcal{R}_F(\mathbb{M}', \{p_2\})$$

408 The above property implies that concretizing any n propositional fluents
 409 cannot resolve foils that weren't resolved by the individual fluents. The above
 410 property follows from the fact that adding a proposition into the model only

411 resolves a foil if it adds a precondition not supported by previous actions in the
 412 plan. Since this is independent of other fluents already part of the abstraction,
 413 we can see that a set of fluents will only resolve the foils that are resolved by
 414 the individual elements of that set.

Proposition 5. *For two models \mathcal{M}_1 , \mathcal{M}_2 and a set of foils F , if $\mathcal{M}_1 = f_{\Lambda}^{safe}(\mathcal{M}_2, \{p_1, \dots, p_k\})$ then for any proposition p ,*

$$\mathcal{R}_F(\{\mathcal{M}_1\}, \{p\}) \supseteq \mathcal{R}_F(\{\mathcal{M}_2\}, \{p\}) \setminus \mathcal{R}_F(\{\mathcal{M}_2\}, \{\})$$

The proposition can be established by following the definition of resolution set and rewriting the lefthand side of the equation as

$$\mathcal{R}_F(\{\mathcal{M}_2\}, p) = \mathcal{R}_F(\{\mathcal{M}_1\}, \{p_1, \dots, p_k\} \cup \{p\})$$

From Proposition 4 we know

$$\begin{aligned} \mathcal{R}_F(\{\mathcal{M}_2\}, p) &= \mathcal{R}_F(\{\mathcal{M}_1\}, \langle p_1, \dots, p_k \rangle \cdot \langle p \rangle) \\ &= \mathcal{R}_F(\{\mathcal{M}_1\}, p) \cup \mathcal{R}_F(\{\mathcal{M}_1\}, \{p_1, \dots, p_k\}) \end{aligned}$$

$$\begin{aligned} \mathcal{R}_F(\{\mathcal{M}_2\}, p) &= \mathcal{R}_F(\{\mathcal{M}_1\}, \langle p_1, \dots, p_k \rangle \cdot \langle p \rangle) \\ &= \mathcal{R}_F(\{\mathcal{M}_1\}, p) \cup \mathcal{R}_F(\{\mathcal{M}_2\}, \{\}) \end{aligned}$$

415 Now removing elements $\mathcal{R}_F(\{\mathcal{M}_2\}, \{\})$ from both LHS and RHS we get

$$\mathcal{R}_F(\{\mathcal{M}_2\}, p) \setminus \mathcal{R}_F(\{\mathcal{M}_2\}, \{\}) = \mathcal{R}_F(\{\mathcal{M}_1\}, p) \setminus \mathcal{R}_F(\{\mathcal{M}_2\}, \{\})$$

416 Which proves our original assertion.

417 This proposition directly leads to the following observation.

418 **Proposition 6.** *Let \mathbb{M}_{min}^F be the minimal abstracting set for a foil set F and*
 419 *\mathcal{M}_H be the human model. if every model in \mathbb{M}_{min}^F is formed from \mathcal{M}_H through*
 420 *f_{Λ}^{safe} , then for any fluent set E_{min} that is optimal for \mathbb{M}_{min}^F then E_{min} must be*
 421 *optimal for \mathcal{M}_H .*

422 We can show the validity of the above proposition through contradiction.
 423 To start with from the definition of foils we know, $\mathcal{R}_F(\{\mathcal{M}_H\}, \{\}) = \emptyset$ and thus
 424 $\mathcal{R}_F(\mathbb{M}_{min}^F, \{\}) = \emptyset$. Let us assume there exists an explanatory set F_1 that is
 425 optimal for human model but not optimal for \mathbb{M}_{min}^F . This could only be due to
 426 two possible reasons, i.e., F_1 is not an explanatory set for \mathbb{M}_{min}^F or there exists
 427 another set F_2 that is optimal for \mathbb{M}_{min}^F but not applicable for \mathcal{M}_H . Through,
 428 Proposition 5 we have already established that any explanatory fluent set for
 429 human model must be an explanatory set for \mathbb{M}_{min}^F . Similarly, from Proposition
 430 1, we know any explanatory set applicable for an abstract model set must be
 431 applicable for the concrete model as well.

432 Now the question is how to exactly identify E_{min} , one possibility is to per-
 433 form an A* search [16] over the space of possible fluent sets to identify E_{min} .
 434 Each search state consists of the minimal set of abstract models for the hu-
 435 man model given the current explanation prefix. We will stop the search as
 436 soon as we find a state where the foils no longer hold for the current minimal
 437 set. In addition to the systematic search, we can see that the specifics of the
 438 setting also allows us to leverage greedy search (described in Algorithm 1). In
 439 each iteration of this search, the algorithm greedily chooses the proposition that
 440 minimizes $\frac{C_p}{|F' \cap \mathcal{R}_F(\mathbb{M}', p)|}$, where F' is the set of unresolved foils at that iteration
 441 and the search ends when all foils are resolved.

442 **Theorem 3.** *The explanatory fluent set \hat{E} generated by Algorithm 1 for a set of*
 443 *foils F and a lattice $\mathcal{L} = \langle \mathbb{M}, \mathbb{E}, \mathbb{P}, \ell \rangle$ is less than or equal to $(\ln k) * C_{E_{min}}^{\mathcal{E}}$, where*
 444 *$C_{E_{min}}^{\mathcal{E}}$ is the cost of an optimal explanatory fluent set and k represents the max-*
 445 *imum number of foils that can be resolved by concretizing a single proposition,*
 446 *i.e., $k = \max_p |\mathcal{R}_F(\mathbb{M}_{min}, p)|$.*

447 *Proof (Sketch).* We will prove the above theorem by showing that Algorithm 1
 448 corresponds to the greedy search algorithm for a weighted set cover problem.
 449 Consider a weighted set cover problem $\langle U, S, W \rangle$ such that the universe set $U =$
 450 F , the subcollections set S is defined as $S = \{s_p | p \in \mathbb{P}\}$ where $s_p = \mathcal{R}_F(\mathbb{M}_{min}, p)$
 451 and the cost of each subset s_p is gives as $W(s_p) = C_p^{\mathcal{E}}$. Proposition 4 ensures

Algorithm 1 Greedy Algorithm for Generating \hat{E}

```

1: procedure GREEDY-EXP-SEARCH
2:   Input:  $\langle F, \mathcal{L} = \langle \mathbb{M}, \mathbb{E}, \mathbb{P}, \ell \rangle \rangle$ 
3:   Output: Explanation  $\hat{E}$ 
4:   Procedure:
5:   curr_model =  $\langle \mathbb{M}_{min}, F \rangle$ 
6:    $\hat{E} = \{\}$ 
7:    $\mathbb{M}_{min} \leftarrow \text{MinimalAbstractModels}(\mathcal{L}, F)$ 
8:   Precompute the resolution sets  $\mathcal{R}_F(\mathbb{M}_{min}, p)$  for each  $p \in \mathbb{P}$ 
9:   while True do
10:     $\mathbb{M}', F' = \text{curr\_model}$ 
11:    if  $|F'| = 0$  then return  $\hat{E}$      $\triangleright$  Return  $\hat{E}$  if all the foils are resolved
12:    else
13:       $p_{next} = \underset{p}{\text{argmin}}(\frac{C_p}{|F' \cap \mathcal{R}_F(\mathbb{M}', p)|})$ 
14:       $\mathbb{M}_{new} = \{\gamma_{p_{next}}(\mathcal{M}) \mid \mathcal{M} \in \mathbb{M}'\}$ 
15:      curr_model =  $\langle \mathbb{M}_{new}, F' \setminus \mathcal{R}_F(\mathbb{M}', p) \rangle$ 
16:       $\hat{E} = \hat{E} \cup p$ 

```

452 that the size of resolution set is a submodular and monotonic function. In
 453 this setting, the act of identifying a set of propositions that resolve the foil
 454 set is identical to coming up with a set cover for U in the new weighted set
 455 cover problem. Furthermore, we can show that the optimal set cover \mathcal{C}_{opt} must
 456 correspond to the cheapest explanation E_{min} (We can prove this equivalence
 457 using Propositions 1,3 and 4, we are skipping the details of this proof due
 458 to space constraints). Algorithm 1 describes a greedy way of identifying the
 459 cheapest set cover for this weighted set cover problem and thus the minimal
 460 explanation for the original problem. For weighted set cover the above greedy
 461 algorithm is guaranteed to generate solutions that are at most $\ln k * W(\mathcal{C}_{opt})$
 462 [17], where $k = \max_{s \in S} |s|$ and this approximation guarantee will hold for E_{min}
 463 as well. \square

464 We can use this algorithm to either generate solutions and or to calculate an
 465 inadmissible heuristic for the previously mentioned A* search. For the heuristic
 466 generation, we will further simplify the calculations (specifically step 8 in Algo-
 467 rithm 1) by considering an over-approximation of \mathcal{R}_F . Instead of considering
 468 the set of all foils resolved by concretizing each proposition p , we will consider
 469 the set of foils where p appears in the precondition of one of the actions in it.
 470 This set should be a superset for \mathcal{R}_F for any proposition.

471 Now that we have formulated the basic form of explanation for this setting,
 472 we will look at how we can relax some of the assumptions made in earlier sections
 473 and how it effects the explanation generation problem. In particular, we will
 474 look at cases where the lattices are no longer proposition conserving, the users
 475 may be raising foils that are sub-optimal as opposed to invalid and finally how
 476 to support models with noise.

477 4.1. Supporting Explanation Generation for Non-Proposition Conserving Lat- 478 tices

479 Proposition conserving lattices, in particular, complete lattices provide a
 480 concise way for the problem designer to specify their knowledge about the end

481 users. In fact, with well-defined abstraction functions, they need only specify
 482 the most concrete model and the set of most abstract models to generate the
 483 rest of the lattice. Unfortunately, there may be cases where such lattices may no
 484 longer be enough to capture all information the system designer may be capable
 485 of providing about the end users. For example, consider a scenario where a robot
 486 needs to put away groceries. The goal of the robot here is to put away a set of
 487 items in prespecified storage locations. In this case, medicines need to be put
 488 in the medicine cabinet while condiments should be placed in kitchen shelves.
 489 In addition to these task-level constraints, the robot’s operations are restricted
 490 by various motion level constraints that limit the possible physical movements
 491 that the robot can perform, including possible ways an object can be grasped
 492 and areas in the workspace it can reach. Clearly, these two types of constraints
 493 are quite different in terms of the background knowledge needed to understand
 494 them. While the task constraints correspond to some simple rules of the task
 495 that are easy to explain to a lay user, understanding the motion constraints
 496 require knowledge about robotics that is usually absent in most users. Thus
 497 there is a natural hierarchy in the concepts related to this task. One way to
 498 capture such information could be by controlling the order in which the various
 499 fluents are considered for abstraction, i.e., remove a particular set of fluents be-
 500 fore moving to others (thereby making the lattice non-proposition conserving).
 501 This means, the easier to understand fluents would get introduced higher up
 502 in the lattice and the harder to understand fluent appear lower in the lattice
 503 closer to the concrete node. The task mentioned above is a particularly good
 504 fit for non-proposition conserving lattices because even the motion constraints
 505 could be captured at multiple conceptual levels. In general, non-proposition
 506 conserving lattices are a useful tool to use when you have settings where there
 507 are different propositions that capture the same phenomena but at varying lev-
 508 els of detail or focus on different aspects. For example, in the case of picking
 509 up an object, one could talk about the ability to pick up the object, picking up
 510 the object by grasping a particular region and even grasping using a particular
 511 grasp point on the object. We can organize the lattice in such a way that the

512 propositions are visited in the order that reflects the preferences of the end-user.
 513 For example, for this scenario, we can arrange the concepts in such a way that
 514 simpler concepts (for example propositions related to simple reachability) are
 515 tested before moving onto more complex concepts.

516 While there are reasons to choose non-proposition conserving lattices and we
 517 could generate explanations using such lattices with some minor modifications
 518 on the solution method described before, the use of such lattices also have a
 519 few disadvantages. The obvious one being that the designer now have to fully
 520 specify such lattices, also the use of such lattices prevents the use of heuristics
 521 and greedy search described in earlier sections. It should also be noted that
 522 when the foils can only be resolved by introducing fluents from lower levels then
 523 the search would still need to search through all the nodes in the above before
 524 identifying the nodes that resolve the foil. Also once such a node is identified, it
 525 won't be easy to separate the set of fluent that actually contribute to resolution
 526 from those that are redundant (particularly when there are multiple foils).

527 To overcome these shortcomings, we will allow designers to specify a non-
 528 proposition conserving lattice while the explanation generation algorithm itself
 529 operates on a modified proposition conserving lattice that uses an updated cost
 530 function. To achieve this, we will start by defining the concept of a well-formed
 531 lattice

532 **Definition 9.** *An abstraction lattice $\mathcal{L} = \langle \mathbb{M}, \mathbb{E}, \mathbb{P}, \ell \rangle$ is said to be **well formed**,
 533 if there exists a unique minimal node (i.e the most concrete model), thus for any
 534 model $\mathcal{M} \in \mathbb{M}$, $\mathcal{M}^\# \sqsubseteq \mathcal{M}$.*

535 Any lattice we describe hence forth, will be assumed to be well-formed unless
 536 specified otherwise. While the concept of minimum abstraction set remains the
 537 same for a non-proposition conserving lattice, analyzing the results of concretiz-
 538 ing the human model with respect to explanatory fluents requires us to look at
 539 a new concept named a completion of a lattice.

540 **Definition 10.** *For a given well formed non-proposition conserving abstraction
 541 lattice $\mathcal{L} = \langle \mathbb{M}, \mathbb{E}, \mathbb{P}, \ell \rangle$, a second lattice $\hat{\mathcal{L}} = \langle \hat{\mathbb{M}}, \hat{\mathbb{E}}, \mathbb{P}, \hat{\ell} \rangle$ is said to be a com-*

542 pletion if $\hat{\mathcal{L}}$ is a proposition conserving lattice, such that, $\mathbb{M} \subseteq \hat{\mathbb{M}}$, $\mathbb{E} \subseteq \hat{\mathbb{E}}$ and
 543 $\ell \subseteq \hat{\ell}$

544 A completion is relevant in this setting, because if we allow the system to
 545 freely choose propositions for the explanatory set, the updated human model (i.e
 546 the model obtained after the explanation) may not be part of the original non-
 547 proposition conserving lattice but is guaranteed to be part of the completion.
 548 Note that completions for a non-proposition conserving lattices are not unique,
 549 but in most cases we will consider a minimal completion. We can create such a
 550 completion by starting with the given lattice and adding any missing incoming
 551 edges iteratively (introducing new models only if there exists no current nodes
 552 that correspond to the set of missing propositions expected at the source of the
 553 edge).

Definition 11. Given a non-proposition conserving lattice $\mathcal{L} = \langle \mathbb{M}, \mathbb{E}, \mathbb{P}, \ell \rangle$, it's
 completion $\hat{\mathcal{L}} = \langle \hat{\mathbb{M}}, \hat{\mathbb{E}}, \mathbb{P}, \hat{\ell} \rangle$, the human model $\mathcal{M}_H \in \mathbb{M}$ and the foil set F , a
 set of propositions $E = \{p_1, \dots, p_n\}$ is said to be a set of explanatory fluents if

$$\forall \pi \in F, \pi(I_{\gamma_E(\mathcal{M}_H)}) \not\models_{\gamma_E(\mathcal{M}_H)} G_{\gamma_E(\mathcal{M}_H)} \text{ and } \gamma_E(\mathcal{M}_H) \in \hat{\mathbb{M}}$$

554 As the original human model is assumed to be part of the given lattice, it
 555 must be part of the completion as well, moreover, the relation between the min
 556 abstraction set and the human model is conserved in the completion as well.
 557 This means that any set of explanatory fluents identified by using the minimum
 558 completion of the given lattice would also be valid for the human model as well.
 559 Such a minimal completion lattice, need not be created beforehand, but could
 560 in fact be generated online when searching for the explanation. Unfortunately,
 561 directly using such a completion lattice for explanation generation (once the min
 562 abstraction set is found), would result in finding sets of propositions that ignore
 563 the information captured by the given lattice. To incorporate this information
 564 we need to not only use the completion we need to consider a new cost function
 565 $C_{\mathcal{L}}^{\mathcal{E}}$ for the explanation generation.

566 **Proposition 7.** Given a min abstraction set \mathbb{M}_{min} for a non-proposition con-

567 serving lattice $\mathcal{L} = \langle \mathbb{M}, \mathbb{E}, \mathbb{P}, \ell \rangle$, we can use it's completion $\hat{\mathcal{L}} = \langle \hat{\mathbb{M}}, \hat{\mathbb{E}}, \mathbb{P}, \hat{\ell} \rangle$ to
 568 identify the explanatory fluents provided the cost of explaining a given proposi-
 569 tion p is defined as $C_{\mathcal{L}}^{\mathcal{E}}(p) = C_p^{\mathcal{E}} + \max_{\mathcal{M} \in \mathbb{M}_{min}} L(p, \mathcal{M})$, where $L(p, \mathcal{M})$ is a penalty,
 570 such that $L(p, \mathcal{M}) \propto C_{\hat{P}}^{\mathcal{E}}$ where \hat{P} is the least costly set of propositions such that
 571 $p \in \hat{P}$ and $\gamma_{\hat{P}}(\mathcal{M}) \in \mathbb{M}$.

572 This new penalty term ensures that a proposition is considered for expla-
 573 nation only after the propositions from higher levels of the given lattice is con-
 574 sidered. Now that we are dealing with explanations using a new proposition
 575 conserving lattice, all earlier results directly carry over including the heuristic,
 576 though the search is less efficient as calculating the cost for each node requires
 577 lookup of the given lattice. Since the proposition conserving lattice is assumed
 578 to be provided upfront, we may be able to precompute the costs.

579 4.2. Supporting Explanations for Sub-optimal Foils

580 We will now consider scenarios where the explainee raises foils that are valid
 581 but may, in fact, be costlier than the one chosen by the robot. In such scenarios,
 582 we would want the robot to explain why the current plan may be preferred, but
 583 such explanations could be complicated by the fact that the actions in the
 584 domain may have state-dependent costs, for example, the cost of picking up a
 585 light block may be lower than picking up a heavier block. Here we would again
 586 need to present the user with a set of fluents and associated action costs that
 587 allow the user to correctly evaluate their alternate plans.

588 To investigate this setting, we will restrict our attention to cases where each
 589 action could be associated a set of positive conditional costs. We will consider
 590 a slightly updated action definition, where each action a for a model \mathcal{M} is now
 591 defined by a tuple of the form $\langle prec_a, e_a^+, e_a^-, \mathcal{C}_a^{\mathcal{M}} \rangle$, where $prec_a, e_a^+$ and e_a^- are
 592 same as before and \mathcal{C}_a are the set of state dependent costs associated with the
 593 action a . $\mathcal{C}_a^{\mathcal{M}}$ is itself defined as a set of individual costs of the form $\langle \phi, c \rangle$,
 594 where ϕ is a conjunction of state literals, which when satisfied in a state causes
 595 the action a to induce a cost c (where $c \in \mathbb{R}_{\geq 0}$). Now the cost of executing the
 596 action a at state s is defined

$$\mathcal{C}_a^{\mathcal{M}}(s) = \sum_{\langle \phi_i, c_i \rangle \in \mathcal{C}_a^{\mathcal{M}}} (\delta(\phi_i, s, c_i))$$

597 Where $\delta(\phi_i, s, c_i) = c_i$ if $s \models \phi_i$ else $\delta(\phi_i, s, c_i) = 0$.

598 We will use the function $\mathcal{C}^{\mathcal{M}}$ to return the total cost of a plan for a given
 599 initial state, i.e, for a plan $\pi = \langle a_1, \dots, a_n \rangle$ and an initial state I , $\mathcal{C}^{\mathcal{M}}(\pi, I) =$
 600 $\mathcal{C}_{a_1}^{\mathcal{M}}(I) + \dots + \mathcal{C}_{a_n}^{\mathcal{M}}(a_{n-1}(\dots(a_1(I))\dots))$.

601 Following the convention set by [18], we can assert that such a domain model
 602 induces a transition system of the form $\mathcal{T} = \langle S, s_0, S_g, L, T, \mathcal{C}^T \rangle$, which is similar
 603 to the original transition system definition except that now each transition is
 604 associated with a cost determined by both source state and action. An abstract
 605 model \mathcal{M}' with a transition system \mathcal{T}' for a set of propositions Λ is defined in
 606 a similar way with the cost of each transition (s, a, s') given by $\mathcal{C}^{\mathcal{T}'}(s, a, s') =$
 607 $\min(\{\mathcal{C}^T(\hat{s}, a, \hat{s}') \mid \hat{s}, \hat{s}' \in S \wedge f_{\Lambda}(\hat{s}) = s \wedge f_{\Lambda}(\hat{s}') = s'\})$.

608 We will also update the explanatory setting a bit and assume that the robot
 609 presents the user with the plan and the anticipated cost of the plan in the most
 610 concrete model (denoted as \mathcal{C}_{π_R}). The user responds by providing a foil set
 611 which they believe is less costlier than the plan in question. Here we can define
 612 a set of explanatory fluents to be

Definition 12. A set of proposition $E = \{p_1, \dots, p_n\}$ is said to be **explanatory fluents** for the human model \mathcal{M}_H and a foil set F if

$$\forall \pi \in F, \pi(I_{\gamma_E(\mathcal{M}_H)}) \not\models_{\gamma_E(\mathcal{M}_H)} G_{\gamma_E(\mathcal{M}_H)} \vee \mathcal{C}^{\mathcal{M}_H}(\pi, I_{\gamma_E(\mathcal{M}_H)}) > \mathcal{C}_{\pi_R}$$

613 Revisiting the abstraction lattice, given the fact that we are dealing with
 614 only positive costs, the first property we can assert is that

615 **Proposition 8.** Given two models \mathcal{M}_1 and \mathcal{M}_2 , such that $\mathcal{M}_1 \sqsubseteq \mathcal{M}_2$, then
 616 for any plan π , we have $\mathcal{C}(\pi, I_{\gamma_E(\mathcal{M}_1)}) \geq \mathcal{C}(\pi, I_{\gamma_E(\mathcal{M}_2)})$

617 This means that once we establish that a given foil is costlier than robot plan
 618 in a model, then it holds in all models that are more concrete than that one. This
 619 insight allows us to reassert Proposition 1 for this new extended definition of

620 explanation and by extension allows us to use the idea of the minimal abstraction
 621 set in this new setting (Proposition 2 holds here as well).

622 This means that we can more or less directly use the search method discussed
 623 for the in-validity case here directly. Unfortunately, in this setting *the size of*
 624 *resolution set is no longer sub-modular* and hence we can not leverage the greedy
 625 method discussed for the pure invalidity case.

626 4.3. Supporting Explanations in the Presence of Human Models with Incorrect 627 Beliefs

628 An underlying assumption for most of the earlier discussion has been the
 629 fact that the user’s model of the task can be represented as an abstraction
 630 of the robot model, i.e. the user model may be imprecise but not incorrect.
 631 Unfortunately, this is not an assumption that can be met in all scenarios. More
 632 often than not, the user may not only be unaware of certain facts pertaining to
 633 the task but may also hold incorrect beliefs about it. Throughout this section,
 634 we will discuss how approaches discussed in earlier sections can be used to handle
 635 such cases.

636 Formally, let the real (but unknown) user model be \mathcal{M}_H and we assert that
 637 this model is an abstraction of some (again unknown) model $\widehat{\mathcal{M}}_R$ that is defined
 638 over the same set of fluents as \mathcal{M}_R , but may have errors in regards to action
 639 definitions, perceived initial and goal state. Let us assume both $\widehat{\mathcal{M}}_R$ and \mathcal{M}_H
 640 belong to the same class of planning problems as defined in Section 2. Again let
 641 the set of alternate plans raised by the user be F . It is important to note that the
 642 reason the user thinks these foils are valid may no longer be just due to missing
 643 fluents, but could also be due to the user’s incorrect understanding of the task.
 644 This means that foils are not an accurate way of identifying the user’s level of
 645 understanding, but we can still use the foils to figure out the level of abstraction
 646 at which the foils can be refuted. Though in scenarios with such models, we have
 647 to consider a complete lattice that contains all possible fluents (i.e assuming user
 648 could be wrong about the use of any of the fluents), i.e., the lattice we will use
 649 would be $\mathcal{L} = \langle \mathbb{M}, \mathbb{E}, \mathbb{P}, \ell \rangle$, where $\mathbb{P} = P_R$ (defined using f_{Λ}^{safe}). We can now

use the methods described in earlier sections to find a set of explanatory fluents \mathcal{E} that can refute the given set of foils. Once the information regarding the explanatory fluents is provided to the user, irrespective of the other fluents, the user should have a correct understanding of each fluents listed in \mathcal{E} . Let $\mathcal{M}_H + \mathcal{E}$ be the updated human model that contains the correct information about \mathcal{E} . Note that even though \mathcal{M}_H or $\mathcal{M}_H + \mathcal{E}$ may not be part of \mathcal{L} , the abstraction of this updated human model that projects out all propositions absent from \mathcal{E} must be part of the lattice \mathcal{L} , i.e., $f_{P \setminus \mathcal{E}}((\mathcal{M}_H) + \mathcal{E}) \in \gamma_{\mathcal{E}}(\mathbb{M}_{min})$. In this scenario, $\gamma_{\mathcal{E}}(\mathbb{M}_{min})$ will be singleton set and we will represent the only element in this set as \mathcal{M}_{min} . As per the definition of valid explanation, we know that $\mathcal{R}_F(\mathbb{M}_{min}, \mathcal{E}) = \emptyset$ and since $\gamma_{\mathcal{E}}(\mathcal{M}_H) \sqsubseteq \gamma_{\mathcal{E}}(\mathcal{M}_{min})$ and therefore the resolution set for $\gamma_{\mathcal{E}}(\mathcal{M}_H)$ must also be empty.

5. Evaluations

5.1. Empirical Evaluations on Explanation Generation for Invalid Foils

For our empirical evaluation, we wanted to understand how effective our basic approaches were in terms of the conciseness of the explanations produced, the solution computation time and the usefulness of approximation. For the approximation, we were interested in identifying the trade-off between decrease in runtime vs. reduction in solution quality. Since both explanation for incorrect beliefs and non proposition-conserving gets compiled down to finding explanation on proposition-conserving lattices, we didn't perform separate evaluations for those methods. All three explanation methods discussed in this paper (blind, heuristic and greedy) were evaluated on five IPC benchmark domains[14]. All the experiments detailed in this section were run on an Ubuntu workstation with 64G RAM.

For each domain, we selected 30 problems from either available test sets or by using standard problem generators (the problems sizes were selected to reflect the size of previous IPC test problems). The lattice for each problem-domain pair was generated by randomly selecting 50% of domain predicates




Figure 3: The graph compares the performance of greedy set cover against the optimal blind search for $|F| = 4$. It plots the average time saved by the set cover and the average increase in cost of the solution for each domain.

and then generating a fully connected proposition conserving lattice using that set of predicates. Since none of the models contained any conditional effects, we created the abstract models by dropping the propositions to be abstracted from the domain models (which are complete for these domains). The foils were generated by selecting random models from the lattice and creating plans from these models that do not hold in the concrete model. Each search evaluated here, generates the set of proposition whose concretizations can resolve the foils set F . In actual applications, this set of propositions needs to be converted into an explanan (the actual message) by considering how this proposition is used in the robot model. Figure 4 shows the explanation generated by our approach for a problem in Rover domain.

Table 1 presents the results from our empirical evaluation on the IPC domains. The table shows the average cost/size of each explanation along with the time taken to generate them. Note that by size, we refer to the number of predicates that are part of the explanation while the cost reflects the total number of unique model updates induced by that explanation. We attempted explanation generation for foil set sizes of one, two and four per problem.

Our main conclusion is that heuristic search seems to outperform blind

Domain Name	$C_{\mathbb{P}}$	$ \mathbb{P} $	$ F $	Blind Search (Optimal)			Heuristic Search			Greedy Set Cover		
				Cost	Size	Time(S)	Cost	Size	Time(S)	Cost	Size	Time(S)
Barman	84.07	7	1	6.87	1	2.43	6.87	1	2.08	6.87	1	3.61
	84	7	2	8.94	1.22	6.35	8.94	1.22	5.71	9.90	1.39	6.05
	90.7	7	4	17.19	1.77	24.99	17.19	1.77	23.7	18.45	1.97	10.34
Rover	168.66	12	1	3.58	1	7.86	3.58	1	5.22	3.58	1	19.18
	188.83	12	2	6.13	1.48	51.36	6.12	1.48	34.04	6.26	1.52	30.5
	192.83	12	4	10.87	2	203.83	10.87	2	181.87	11.42	2.19	49.32
Satellite	53.01	4	1	18.73	1	2.23	18.73	1	1.92	18.73	1	1.49
	60.77	4	2	32	1.61	7.21	32	1.6	5.86	32.53	1.7	3.04
	62.73	4	4	43.27	2.29	18.67	43.27	2.29	16.42	43.88	2.39	5.85
Woodworking	156.71	7	1	14.45	1	2.84	14.45	1	2.23	14.45	1	3.35
	146.33	7	2	20.62	1.21	6.88	20.62	1.21	4.93	21.38	1.38	6.25
	154	7	4	28.62	1.69	24.70	28.62	1.69	19.49	30.41	2	12.13
Sokoban	220.6	3	1	51.21	1	1.51	51.21	1	1.35	51.21	1	1.28
	151.72	3	2	94.52	1.55	3.93	94.52	1.55	3.35	98.31	1.73	2.59
	220.69	3	4	136.41	2.22	8.75	136.41	2.22	8.3	141.93	2.37	5.23

Table 1: Table showing runtime/cost for explanations generated for standard IPC domains. Column $|\mathbb{P}|$ represents number of predicates that were used in generating the lattice, while $C_{\mathbb{P}}^{\varepsilon}$ represents the cost of an explanation that tries to concretize all propositions in \mathbb{P} and provides an upper bound on explanation cost.

1. Calibrate camera to objective0 2. Take an image of objective0 3. Communicate the image to the lander 4. Communicate the soil data to the lander 5. Communicate the rock data to the lander	Predicate to concretize with: have_soil_analysis Explanation for affected actions: <ul style="list-style-type: none"> • have_soil_analysis is required as a precondition for communicate soil data, but is false at step 4 of the foil • have_soil_analysis is part of the add effects for the sample soil action
Human's Foil	Robot Explanation

Figure 4: An example explanation generated by our system for IPC rover domain. The human incorrectly believes that the rover can communicate sample information without explicitly collecting any samples. While the abstraction lattice in this example was generated by projecting out upto 12 predicates, the search correctly identifies concretizations related to $(have_soil_analysis\ ?r - rover\ ?w - waypoint)$ as the cheapest explanation ($C_E^{\mathcal{E}} = 2$ as opposed to $C_P^{\mathcal{E}} = 55$)

697 search in almost every problem and generates near-optimal solutions (Blind
698 search always generates the minimal explanation). Further, we saw that greedy
699 search outperformed heuristic search in most cases barring a few exceptions.
700 The greedy search was able to make significant gains especially for higher foil
701 set sizes. This is entirely expected due to the fact that step 8 in Algorithm
702 1 can be expensive for problems with long plans (but still polynomial). This
703 expensive pre-computation pays off as we move to cases where E_{min} consists
704 of multiple propositions. Additionally, we found out that greedy solutions were
705 quite comparable to the optimal solutions with respect to their costs. For exam-
706 ple in $|F| = 4$ for satellite domain, while the greedy solution cost took a penalty
707 of $\sim 1.4\%$ the search time was reduced by $\sim 68\%$. Figure 3 plots the compari-
708 son between the time saved by the greedy search versus any loss in optimality
709 incurred by the greedy search.

710 5.2. Empirical Evaluations on Explanation Generation for Sub-optimal Foils

711 Next, we wanted to evaluate the empirical performance of the approach for
712 domains with state dependent cost. For this setting, since we don't have stan-
713 dard benchmark domains with this property, we chose standard IPC domains
714 and modified them to include conditional cost updates. In particular, we chose
715 blocksworld, zenotravel, gripper and rover. For blocksworld, we introduced
716 three new predicates, namely `heavy`, `light` and `unsteady` each of which takes a
717 block as an argument. For each problem instance, we assigned each block to be
718 either heavy or light and set some of the blocks as unsteady. We also updated
719 the stack action so that stacking a heavy block on a light one or an already
720 unsteady one cause the block to be unsteady. We also set a high cost penalty
721 for stacking any block on an unsteady one. For zenotravel, we came up with
722 three binary predicates `near`, `farther` and `farthest` that takes cities as arguments.
723 We also assigned a higher cost for traveling between far away cities than nearby
724 ones (so the optimal plan may involve the plane making a lot more stops). For
725 gripper, we again mark a ball to be `heavy` or `light` and now each robot can also
726 pick up two balls at the same time. We assign a high cost to picking up heavy
727 balls and picking up the second ball in a gripper that is already holding a ball.
728 We also provide the robot with a push action, that allows for it to move heavy
729 balls without accruing large cost. Finally in the case of rover domains, we set
730 some of the waypoints as being hilly area and communicating from these way-
731 points are assigned higher costs. Table 2 presents the explanation generation
732 time and average explanation sizes for the modified domains. For each domain,
733 we generated five problems and the test was run using systems of the same con-
734 figuration as Section 5.1. For Blocksworld we considered instances where the
735 number of blocks spanned from four blocks to 20, for Gripper all problems had
736 two rooms and up to 12 balls and for Rover domain problems had upto three
737 objectives and four waypoints. Finally, in Zenotravel all problems considered
738 traveling between 10 cities and the number of passengers ranged from 20 to
739 60. The fact plan being explained were generated using optimal planners when
740 possible and the foils were generated either using a satisficing planner (Metric-

Domain	Average Explanation Size	Runtime
BLOCKS	4.4	8.319
Gripper	5.4	7.368
Rover	4	9.690
Zeno	6.6	8.905

Table 2: The sample runtime and average explanation size for five problem instances from the modified domains.

FF [19]) or hand written using knowledge about the domain. As expected, the search was able to find the minimal number of predicates to be included into the problem to resolve the foils, for example in Blocksworld, the approach was able to correctly identify the predicate unsteady as being enough to explain the foils in the example.

In addition to the empirical results discussed in this paper on classical planning problems, the approaches discussed have also shown to be useful in modeling explanatory dialogue in the context of Task and Motion Planning (as shown in [20]).

5.3. User Study to Evaluate Role of Abstractions in Explanation

In this section, we will consider one of the assumptions that we made throughout the work, namely that providing the explanations at an abstract level would help reduce the cognitive burden on the user’s end. Specifically, we will test the following hypothesis

Hypothesis 1. *Given two models \mathcal{M}_1 and \mathcal{M}_2 , such that $\mathcal{M}_1 \sqsubset \mathcal{M}_2$ and \mathcal{M}_2 is formed using methods presented in Section 2.1, a user would find it easier to work with the more abstract model \mathcal{M}_2 when compared to \mathcal{M}_1*

We will evaluate this hypothesis over two different dimensions. One with respect to the subjective workload the user may experience when working with such a model to achieve some task, and then with respect to the actual ability of the user to successfully complete the task. For the former, we will employ

762 NASA-Task Load Index (NASA-TLX) survey [21], while for the latter we mea-
763 sure the time taken by the user to complete the task. NASA-TLX is a very
764 influential and widely used method to gauge the subjective workload experi-
765 enced by the user. NASA-TLX, divides the workload of a task over six different
766 dimensions; namely, Mental Demand, Physical Demand, Temporal Demand,
767 Effort, Performance, and Frustration. The users are first asked to rate the task,
768 across these dimensions on a 20 point scale (with larger value denoting higher
769 workload). They are then required to provide relative weights across these di-
770 mensions, by making pairwise choices between these different dimensions. A
771 weighted average of the ratings provided across these dimensions is then used
772 as a measure of the workload.

773 For the actual study, we relied on a between-subject study design, wherein
774 the study participants are divided into two groups. All study participants were
775 students from ASU. One received abstract explanations and the other group was
776 given concrete domain model information as explanations. As the task in ques-
777 tion, we used a variation of the Sokoban domain, that involves the agent pushing
778 a box to a pre-specified domain. Unlike the common versions of Sokoban, this
779 variant involved the robot needing to first turn on a switch before pushing the
780 boxes. Each participant in the study was allowed to play the game through a
781 web-interface (which is shown in Figure 5, along with a sample explanation).
782 While they were told the actions they can perform, they weren’t told what
783 each of the action achieves or their preconditions. Each player was allotted a
784 total time of five minutes to complete the game. As the users play the game
785 and if they perform an invalid action, they were provided with an explanation
786 appropriate for their group.

787 For both groups, the current action sequence being executed was treated as
788 the foil and the explanation consisted of the following information; the state at
789 which the sequence failed, the specific action that failed, the expected set of pre-
790 conditions, the failed precondition, and lifted model information about relevant
791 actions. While one group of users were shown the information with respect to
792 the concrete model, i.e., they were shown the full state, all the preconditions,

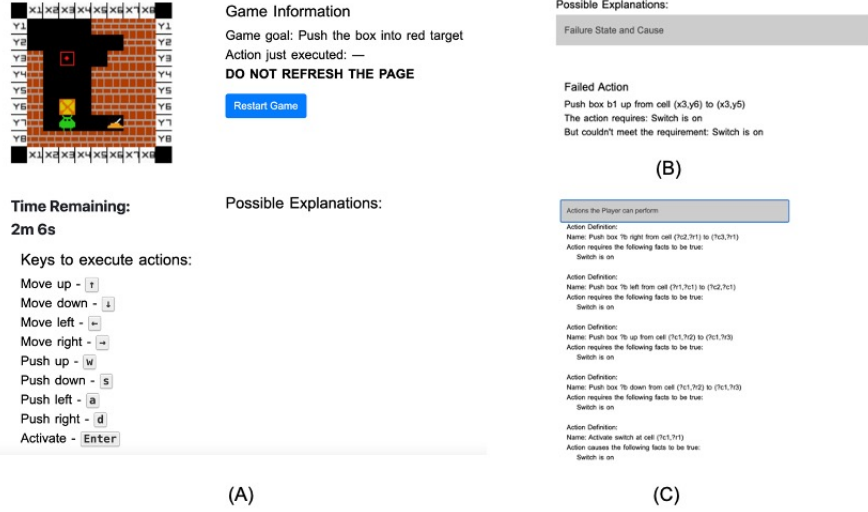


Figure 5: Screenshots from the user interface exposed to the end user. (A) The participant is shown the current state of the game, they are allowed to control the agent via their keyboard and whenever they perform an invalid action, they are shown a possible explanation. (B) and (C) presents a sample explanation provided to participants who were exposed to abstract explanations. Here the current state shown to the participant is empty as none of the facts are true in that state.

all failed preconditions, and the entire model of the task, the second group was shown the abstracted version of each of the above-mentioned information. The level of abstraction for this group is identified based on the foil failure. To make sure the explanation generation time is symmetric between the groups, we avoided search to identify the best level of abstraction and rather we simulated the failing foil in the most concrete model and randomly selected the predicate corresponding to one of the failing precondition to generate the abstract models. While this may result in a more detailed model than required, as we will see, even with this simple approach we did see a significant difference between the two groups. We also carried over predicates from consecutive failures, so the users of the second group saw increasingly more concrete models if they failed repeatedly (though still more abstract compared to the first group).

Scale	Concrete-Explanation Group	Abstract-Explanation Group
Effort	1.595	1.252
Frustration	2.5	2.19
Mental Demand	2.9	1.414
Performance	1.186	1.705
Physical Demand	0.038	0.152
Temporal Demand	1.929	1.719

Table 3: The weighted average workload reported by the participants of the user study across the individual scales used in NASA-TLX.

805 In total, we collected responses from 28 participants, 14 of whom had access
 806 to concrete explanation (henceforth referred to as Concrete-explanation group),
 807 and the remaining 14 were provided with abstract explanation (i.e the Abstract-
 808 explanation group). While the Concrete-explanation group on average took
 809 200.857 secs to finish the task, the abstraction group only took 163.5 seconds.
 810 In terms of the weighted average workload for the Concrete-explanation group,
 811 we saw 10.147 and for the Abstract-explanation group, we found it to be 8.433.
 812 The distribution across the six scales are presented in Table 3. As seen from the
 813 table, in all but Performance and Physical demand, people reported a higher
 814 workload for the concrete explanation group. We see a particularly significant
 815 difference across the mental demand dimension, which was the main focus of
 816 the assumptions made by our work. Thus the results from both the subjective
 817 workload study, and the performance of the user (measure in terms of the time
 818 taken by the user to finish the task), conform to our original hypothesis and
 819 we see that the use of abstractions provides a distinct advantage over providing
 820 complete details.

821 6. Related Work

822 There is increasing interest within the automated planning community to
 823 solve the problem of generating explanations for plans ([22, 23]). Earlier works

like [3, 4, 24] looked at explanations as a way of describing the effects of plans, while works like [25, 26] looked at plans itself as explanations for a set of observations. Another approach that has received a lot of interest recently is to view explanations as a way of achieving model reconciliation [27]. Such explanations are seen as a solution to a *model reconciliation problem* (referred to as MRP) and this approach postulates that the goal of an explanation is to update the model of the observer so they can correctly evaluate the plans in question. The methods discussed in this paper can be seen as performing a type of model-reconciliation, but one could also leverage the methods discussed here to relax some of the assumptions made by model-reconciliation works for certain conditions. We discuss the relationship between model-reconciliation and the methods studied in this paper in more detail in Section 7.1.

As noted, our work is closely related to the well studied method of counter-example guided refinement or CEGAR that was originally developed for Model checking [5]. Many planning works have successfully used CEGAR based methods to generate heuristics for plan generation [8, 28]. The idea of foil resolution set for a given concretization is also closely related to the process of identifying spurious counter examples employed by CEGAR based methods (cf. [29, 9, 30]). One major difference between our work and standard CEGAR based methods is the fact that in our setting the abstract model producing the foil (or counter-example) is unknown. Since we are exclusively dealing with spurious counter-examples we are also not bound to testing our foils (in other words identifying faults or pivot states) in the most concrete model (which could be quite expensive). Further, traditional CEGAR methods are generally not as focused on identifying the cheapest refinements.

Many abstraction schemes have been proposed for planning tasks (starting with [7]), but in this paper, we mainly focused on state abstractions and based our formulation on previous works like [13] and [12]. It would be interesting to see how we can extend the approaches discussed in this paper to handle temporal and procedural abstractions (e.g., HLAs [31]).

There exists a rich body of literature that has debated and discussed the role

855 of abstraction in Social sciences (cf. [32, 33] for arguments towards abstraction,
 856 while [34] argues for adding more details provided the task constraints allow
 857 for it). Unlike these works that study explanation in everyday scenarios, expla-
 858 nation in the context of AI systems have a markedly different flavor, in so far
 859 that the explainer may be representing and reasoning about the task at levels of
 860 details that may be too hard for the users to understand. Thus abstraction can
 861 be a powerful tool in identifying just the required level of information to allow
 862 people to achieve their goals. This is an intuition being leveraged by more and
 863 more works to help generate explanations or even decisions that are easier to
 864 understand. For example, state abstractions have been leveraged by [35] to gen-
 865 erate simpler models that generate easier to understand policies, and [36] uses
 866 abstraction to simplify policies. Even in the realm of machine learning explana-
 867 tions, abstractions have been considered as a way to generate multi-resolution
 868 explanations [37]. The importance of adjusting the level of details for different
 869 users have also been considered and argued in [38], where they propose three
 870 levels of explanations, namely, high-level, low-level, and co-created level expla-
 871 nations. While high and low-level explanations focus on generating summaries
 872 and detailed descriptions respectively, co-created explanations use the user in-
 873 teraction to determine the contents of the explanation. Our specific methods
 874 could be considered closely related to the co-created explanation studied in the
 875 paper.

876 There have also been recent works that have looked at generating contrastive
 877 explanations for planning. Some significant examples for these include works
 878 like [39] and [40]. Both these cases treat the cause of user’s confusion to be
 879 their limited computational capabilities and the explanations tend to help them
 880 realize the consequences of following the foil without worrying about model
 881 reconciliation.

882 A closely related but distinct form of explanations is the one where the
 883 explanan (i.e. the information provided to the explainee) constitutes a counter-
 884 factual example [41]. Such explanations are particularly popular in classification
 885 settings, where when queried about an inexplicable classification, the system re-

sponds with a counterfactual example where the desired decision may have been made. Note that in such cases, the system needs to focus on generating counterfactual instances that the user would find acceptable. Many recent works have looked at identifying desirable properties for such counterfactual explanations (cf. [41, 42]), and some of the prominent ones identified in the literature include, making sure the counterfactual example is close enough to the decision-point in question and the counterfactual is plausible, in terms of not only being a plausible datapoint but also that it is actionable. Actionability can be particularly important in domains like loan approval, wherein the counterfactual represents the changes the user needs to make to achieve the desired outcomes. Note that in our method, it is the user who is responsible for generating the counterfactual example and as such is guaranteed to come up with foil they believed to be most likely or most useful. Thus our focus has been on ensuring that the explanations generated in response meet the desired properties discussed in the literature. As discussed above, our explanations do meet many of the important requirements discussed in the literature including being selective and social.

7. Conclusion and Discussion

In this paper, we investigated the problem of generating explanations when the explainee understands the task model at a higher level of abstraction. We looked at how we can use explanations as concretization for such scenarios and proposed algorithms for generating minimal explanations. One unique aspect of our approach is the use of foils as a way of capturing human confusion about the problem. This not only helps us formulate more efficient explanation generation methods but also aligns with the widely held belief that human expect contrastive explanations (cf. [43, 44]). Moreover, in most real-world scenarios humans usually include the foil in the request for explanations unless the foil is quite apparent from the context. The use of abstractions in explanations also allows us to reduce the cognitive burden imposed on the user for understanding the explanations. Below we have provided some more detailed discussion on

915 the nature of explanation generated by the methods discussed in the paper and
916 some future work.

917 7.1. HELM and Model Reconciliation Explanation

918 As mentioned earlier, the methods discussed in this case could be seen as a
919 special case of model reconciliation [27]. Here the model updates are limited to
920 model concretization and the human’s model is an abstraction of the original
921 model. Rather than assuming that we are given an explicit human model,
922 we assume that the human model belongs in the set of possible models that
923 corresponds to the various abstractions of the robot model. In this sense, this
924 method is also comparable to the work done on generating explanations for a set
925 of possible models [45], and in particular to the conformant explanations studied
926 in that paper. Though unlike [45], in this setting we can guarantee that the
927 conformant explanation is also minimal for the unknown human model (provided
928 all model updates and hence explanations are restricted to model concretizations
929 over fluents). Our use of minimal abstractions for explanations also allows the
930 methods to handle cases where the user questions arise due to a mismatch in
931 the inferential capabilities and not just a mismatch in the knowledge about the
932 task. While the original model reconciliation work focused on explanations that
933 address all possible foils, our work specifically tries to address foils raised by the
934 user. This allows us to provide more concise explanations and allows us to scale
935 to larger problems as compared to the original MRP approaches.

936 Another way to connect this work with model reconciliation is to leverage
937 the insights from Section 4.3, to show that the method described in this pa-
938 per can also be used in the more general model-reconciliation setting. Section
939 4.3 shows that for the class of planning problems studied in this paper, even
940 when the human model may not meet the assumption that it is, in fact, an ab-
941 straction of the original robot model, we can still generate an explanation that
942 refutes the given set of foil using abstractions formed from the robot model.
943 Provided we use a complete abstraction lattice which contains a single maximal
944 node formed by projecting out all the propositions. This means for explana-

945 tory queries related to just refuting alternate plans, abstraction lattices give us
 946 a way to circumvent one of the most restrictive assumptions made in model-
 947 reconciliation works, namely the need to know or learn the human model. As
 948 discussed in Section 4.3, the explanations generated over abstraction lattices
 949 will remain valid model-reconciliation explanations regardless of how the hu-
 950 man model may be different from the robot model (provided it is still of the
 951 form described in Section 2 and doesn't contain any fluents absent from the
 952 robot model). Though compared to model-reconciliation techniques like those
 953 studied in [27, 45], the methods discussed in this paper could generate much
 954 larger explanations. For one, the explanations here involve providing informa-
 955 tion about all the uses of the explanatory fluents in the robot model, many
 956 of which the user may already know. This approach can also be extended to
 957 generate explanations of unsolvability and for partial foils of the type discussed
 958 in [46].

959 7.2. Properties of Explanations

960 The prior work on explanation as model reconciliation [27] mainly used four
 961 properties to characterize the various types of explanations that were introduced
 962 in the paper. These properties were Completeness, Monotonicity, Conciseness,
 963 and Computability. We too can use these properties to describe the explanations
 964 we have looked at (with small updates to meet our specific setting).

965 Any explanation generated by our methods will be complete and monotonic.
 966 While [27] defines a complete explanation as one that guarantees optimality
 967 of the plan under question. For our scenario, a complete explanation can be
 968 redefined as one that resolved all the given foils ($|\mathcal{R}_F(\Pi', E)| = |F|$). [27]
 969 considers an explanation monotonic if no future explanation can invalidate it.
 970 In our setting, this means that once a foil has been resolved by an explanation,
 971 no future explanation (or model concretizations) can reintroduce it. Which is
 972 satisfied by any explanation as concretization.

973 As for the remaining two properties (Conciseness and Computability), the
 974 definitions laid out in the original MRP paper directly applies for our setting

as well. Similar to MRP explanations, computability and conciseness remains incompatible properties for explanations in our case too. The explanations that are easier to compute end up being neither concise nor easy to understand. For example, one simple strategy to provide explaining a plan would be to provide enough details to the explainee that the human model completely converges to the robot model, but this strategy could be extremely expensive and even unnecessary.

In addition to properties discussed in [47], works in social sciences have also prescribed some essential characteristics for what would be considered useful explanations by people [2]. Chief among them is generating contrastive explanations, which remains the central thrust of the methods discussed in this work. The other two properties usually cited by such sources are *selectiveness* and being *social*. An explanation is considered selective if it chooses to focus only on the aspects relevant to the current explanatory query. As such this is directly related to the minimality of explanation and thus the methods discussed in this paper can be considered to be selective. On the other hand, an explanation is considered social if it is tailored to the user’s background. Our method supports this property in two distinct ways, one by explicitly trying to localize the user’s model on the abstraction lattice, and by allowing the abstraction lattice itself to be tailored to reflect the preferences of the users.

7.3. Other Explanatory Queries

The explanatory approaches discussed in this paper have mostly focused on helping users resolve their confusion about foils, but it may also be possible that they have questions about the original robot plan. For the robot plan $\pi_R = \langle a_1, \dots, a_n \rangle$, user could raise questions of the following types

1. Why perform action a_i ? (where $a_i \in \pi_R$)
2. How can action a_i be performed when the precondition p is not met?
(where $a_i \in \pi_R$ and $p \in \text{prec}_{a_i}^+$ or $p \in \text{prec}_{a_i}^-$ in the human model \mathcal{M}_H)

Question (1) captures the user’s concerns regarding the use of any particular action in the plan, while question (2) captures their concerns regarding the

1005 validity of the plan. Other questions, such as achievement of goals and questions
1006 about the overall plan can be cast in terms of these more basic questions. For
1007 answering questions of the first type, we can easily adapt approaches discussed
1008 in works like [3]. For a given action, these approaches try to find causal links
1009 that capture the specific action’s contributions. We can leverage the hierarchy
1010 specified by the abstraction lattice to identify causal links consisting of higher-
1011 level concepts.

1012 For Question (2), it is possible to view such questions as another type of foil.
1013 While in earlier sections we tried to find abstract models where a particular foil
1014 can be refuted, here we just need to find the level at which the specified pre-
1015 conditions can be met. In the absence of disjunctive preconditions, we wouldn’t
1016 need to perform a search to find such models, but rather choose the first ab-
1017 stract model where fluents corresponding to the preconditions in question is
1018 introduced.

1019 This paper mostly focuses on cases where foils are fully specified. Such fully
1020 specified foils may not always be available and the human may instead be only
1021 ready to specify certain parts of the foil. In such cases, the exact foil set would
1022 consist of all plans that could potentially satisfy the plan level constraints being
1023 specified by the user question. In [46], the methods discussed in this paper have
1024 been extended to handle such cases. The work tries to handle such partial foil
1025 specifications by compiling it directly into each model in the abstraction lattice
1026 without generating the complete set of foils. Though the work only looks at
1027 employing blind search to generate such explanations. Such abstraction based
1028 explanations have also been used to generate explanations in the context of
1029 providing assistance for domain-authoring tools [48].

1030 **Acknowledgement.** This research is supported in part by ONR grants N00014-
1031 16-1-2892, N00014-18-1-2442, N00014-18-1-2840, N00014-9-1-2119, AFOSR grant
1032 FA9550-18-1-0067, DARPA SAIL-ON grant W911NF-19-2-0006, NSF grants
1033 1936997 (C-ACCEL), 1844325 and 1909370, NASA grant NNX17AD06G, and
1034 a JP Morgan AI Faculty Research grant.

References

- [1] T. Chakraborti, S. Sreedharan, S. Kambhampati, The emerging landscape of explainable automated planning & decision making, in: Proceedings of the Twenty-Ninth International Joint Conference on Artificial Intelligence, IJCAI-20, 2020, pp. 4803–4811.
- [2] T. Miller, Explanation in artificial intelligence: Insights from the social sciences, CoRR abs/1706.07269 (2017). URL: <http://arxiv.org/abs/1706.07269>.
- [3] B. Seegebarth, F. Müller, B. Schatttenberg, S. Biundo, Making hybrid plans more clear to human users-a formal approach for generating sound explanations, in: Twenty-Second International Conference on Automated Planning and Scheduling, 2012.
- [4] P. Bercher, S. Biundo, T. Geier, T. Hoernle, F. Nothdurft, F. Richter, B. Schatttenberg, Plan, repair, execute, explain-how planning helps to assemble your home theater., in: ICAPS, 2014.
- [5] E. Clarke, O. Grumberg, S. Jha, Y. Lu, H. Veith, Counterexample-guided abstraction refinement, in: International Conference on Computer Aided Verification, Springer, 2000, pp. 154–169.
- [6] S. Sreedharan, S. Srivastava, S. Kambhampati, Hierarchical expertise-level modeling for user specific contrastive explanations, in: IJCAI, 2018.
- [7] E. D. Sacerdoti, Planning in a hierarchy of abstraction spaces, Artificial intelligence 5 (1974) 115–135.
- [8] J. Seipp, M. Helmert, Counterexample-guided cartesian abstraction refinement., in: ICAPS, 2013.
- [9] E. R. Keyder, J. Hoffmann, P. Haslum, et al., Semi-relaxed plan heuristics., in: ICAPS, 2012.

- 1061 [10] J. C. Culberson, J. Schaeffer, Pattern databases, *Computational Intelli-*
1062 *gence* 14 (1998) 318–334.
- 1063 [11] S. Edelkamp, Planning with pattern databases, in: *ECP*, 2000.
- 1064 [12] C. Backstrom, P. Jonsson, Bridging the gap between refinement and heuris-
1065 *tics in abstraction*, in: *Twenty-Third International Joint Conference on*
1066 *Artificial Intelligence*, 2013.
- 1067 [13] S. Srivastava, S. J. Russell, A. Pinto, *Metaphysics of planning domain*
1068 *descriptions.*, in: *AAAI*, 2016, pp. 1074–1080.
- 1069 [14] International Planning Competition, IPC Competition Domains, <https://goo.gl/i35bxc>, 2011.
1070
- 1071 [15] K. Bernhard, J. Vygen, *Combinatorial optimization: Theory and algo-*
1072 *rithms*, Springer, Third Edition, 2005. (2008).
- 1073 [16] P. E. Hart, N. J. Nilsson, B. Raphael, *A formal basis for the heuristic de-*
1074 *termination of minimum cost paths*, *IEEE transactions on Systems Science*
1075 *and Cybernetics* 4 (1968) 100–107.
- 1076 [17] N. E. Young, *Greedy set-cover algorithms*, in: *Encyclopedia of algorithms*,
1077 Springer, 2008, pp. 1–99.
- 1078 [18] F. Geißer, T. Keller, R. Mattmüller, *Abstractions for planning with state-*
1079 *dependent action costs*, *ICAPS* (2016).
- 1080 [19] J. Hoffmann, *The metric-ff planning system: Translating“ignoring delete*
1081 *lists”to numeric state variables*, *Journal of artificial intelligence research*
1082 20 (2003) 291–341.
- 1083 [20] S. Sreedharan, M. P. Madhusoodanan, S. Srivastava, S. Kambhampati,
1084 *Plan Explanation Through Search in an Abstract Model Space: Extended*
1085 *Results*, in: *ICAPS XAIP Workshop*, 2018.

- 1086 [21] S. G. Hart, L. E. Staveland, Development of nasa-tlx (task load index):
1087 Results of empirical and theoretical research, in: *Advances in psychology*,
1088 volume 52, Elsevier, 1988, pp. 139–183.
- 1089 [22] M. Fox, D. Long, D. Magazzeni, Explainable Planning, in: *IJCAI XAI*
1090 *Workshop*, 2017.
- 1091 [23] P. Langley, B. Meadows, M. Sridharan, D. Choi, Explainable Agency for
1092 Intelligent Autonomous Systems, in: *AAAI/IAAI*, 2017.
- 1093 [24] S. Kambhampati, A classification of plan modification strategies based on
1094 coverage and information requirements, in: *AAAI 1990 Spring Symposium*
1095 *on Case Based Reasoning*, Citeseer, 1990.
- 1096 [25] S. Sohrabi, J. A. Baier, S. A. McIlraith, Preferred explanations: Theory
1097 and generation via planning., in: *AAAI*, 2011.
- 1098 [26] B. L. Meadows, P. Langley, M. J. Emery, Seeing beyond shadows: Incre-
1099 mental abductive reasoning for plan understanding., in: *AAAI Workshop:*
1100 *Plan, Activity, and Intent Recognition*, volume 13, 2013, p. 13.
- 1101 [27] T. Chakraborti, S. Sreedharan, Y. Zhang, S. Kambhampati, Plan expla-
1102 nations as model reconciliation: Moving beyond explanation as soliloquy,
1103 in: *IJCAI*, 2017.
- 1104 [28] J. Seipp, M. Helmert, Diverse and additive cartesian abstraction heuristics.,
1105 in: *ICAPS*, 2014.
- 1106 [29] P. Haslum, J. Slaney, S. Thiébaux, et al., Incremental lower bounds for
1107 additive cost planning problems., in: *ICAPS*, volume 12, 2012, pp. 74–82.
- 1108 [30] M. Steinmetz, J. Hoffmann, Towards clause-learning state space search:
1109 Learning to recognize dead-ends., in: *AAAI*, 2016, pp. 760–768.
- 1110 [31] B. Marthi, S. J. Russell, J. A. Wolfe, Angelic semantics for high-level
1111 actions., in: *ICAPS*, 2007, pp. 232–239.

- 1112 [32] A. Garfinkel, *Forms of explanation: Rethinking the questions in social*
1113 *theory*, Yale University Press, 1982.
- 1114 [33] C. Hitchcock, J. Woodward, *Explanatory generalizations, part ii: Plumbing*
1115 *explanatory depth*, *Noûs* 37 (2003) 181–199.
- 1116 [34] C. Bechlivanidis, D. A. Lagnado, J. C. Zemla, S. Sloman, *Concreteness*
1117 *and abstraction in everyday explanation*, *Psychonomic bulletin & review*
1118 24 (2017) 1451–1464.
- 1119 [35] J. Ferrer-Mestres, T. G. Dietterich, O. Buffet, I. Chadès, *Solving k-mdps*,
1120 *in: Proceedings of the International Conference on Automated Planning*
1121 *and Scheduling*, volume 30, 2020, pp. 110–118.
- 1122 [36] N. Topin, M. Veloso, *Generation of policy-level explanations for reinforce-*
1123 *ment learning*, *in: Proceedings of the AAAI Conference on Artificial Intel-*
1124 *ligence*, volume 33, 2019, pp. 2514–2521.
- 1125 [37] D. Bayani, S. Mitsch, *Fanoos: Multi-resolution, multi-strength, interactive*
1126 *explanations for learned systems*, *in: IJCAI XAI Workshop*, 2020.
- 1127 [38] K. Martin, A. Liret, N. Wiratunga, G. Owusu, M. Kern, *Developing a*
1128 *catalogue of explainability methods to support expert and non-expert users*,
1129 *in: International Conference on Innovative Techniques and Applications of*
1130 *Artificial Intelligence*, Springer, 2019, pp. 309–324.
- 1131 [39] R. Eifler, M. Cashmore, J. Hoffmann, D. Magazzeni, M. Steinmetz, *A new*
1132 *approach to plan-space explanation: Analyzing plan-property dependencies*
1133 *in oversubscription planning*, AAAI, 2020.
- 1134 [40] M. Cashmore, A. Collins, B. Krarup, S. Krivic, D. Magazzeni,
1135 D. Smith, *Towards explainable ai planning as a service*, *arXiv preprint*
1136 *arXiv:1908.05059* (2019).
- 1137 [41] M. T. Keane, B. Smyth, *Good counterfactuals and where to find them: A*
1138 *case-based technique for generating counterfactuals for explainable ai (xai)*,
1139 *CBR*, 2020.

- 1140 [42] R. M. Byrne, Counterfactuals in explainable artificial intelligence (xai):
1141 Evidence from human reasoning., in: IJCAI, 2019, pp. 6276–6282.
- 1142 [43] T. Lombrozo, Explanation and abductive inference, Oxford handbook of
1143 thinking and reasoning (2012) 260–276.
- 1144 [44] T. Lombrozo, The structure and function of explanations, Trends in Cog-
1145 nitive Sciences 10 (2006) 464 – 470.
- 1146 [45] S. Sreedharan, S. Kambhampati, et al., Handling model uncertainty and
1147 multiplicity in explanations via model reconciliation, in: Proceedings of
1148 the International Conference on Automated Planning and Scheduling, vol-
1149 ume 28, 2018.
- 1150 [46] S. Sreedharan, S. Srivastava, D. Smith, S. Kambhampati, Why can’t you
1151 do that hal? explaining unsolvability of planning tasks, in: Proc. IJCAI,
1152 2019.
- 1153 [47] T. Chakraborti, S. Kambhampati, M. Scheutz, Y. Zhang, AI Challenges in
1154 Human-Robot Cognitive Teaming, arXiv preprint arXiv:1707.04775 (2017).
- 1155 [48] S. Sreedharan, T. Chakraborti, C. Muise, Y. Khazaeni, S. Kambhampati,
1156 D3wa+: A case study of xaip in a model acquisition task, in: Proc. ICAPS,
1157 2020.