

Using Deception in Markov Game to Understand Adversarial Behaviors through a Capture-The-Flag Environment



Siddhant Bhambri^{1*}



Purv Chauhan^{1*}



Frederico Araujo²



Adam Doupé¹



Subbarao Kambhampati¹

Objective

- ❖ Identifying real-world adversarial threats
- ❖ Advantage of deploying deception strategies
- ❖ Modeling the real-world attacker
- ❖ Role of game-theoretic decision models
- ❖ Understanding human attacker behaviors



Agenda

- ❖ Reviewing Honey-Patching
- ❖ Setting up the user-study
- ❖ Deploying Honey-Patches as mitigations
- ❖ Modeling the adversary's attack graph
- ❖ Formulating the Markov Game model
- ❖ Computing the Bayesian Stackelberg Equilibrium

Traditional Solution to Mitigate Attacks



REPORTED!

**DEPLOY
UPDATES/PATCHES**

Traditional Solution to Mitigate Attacks



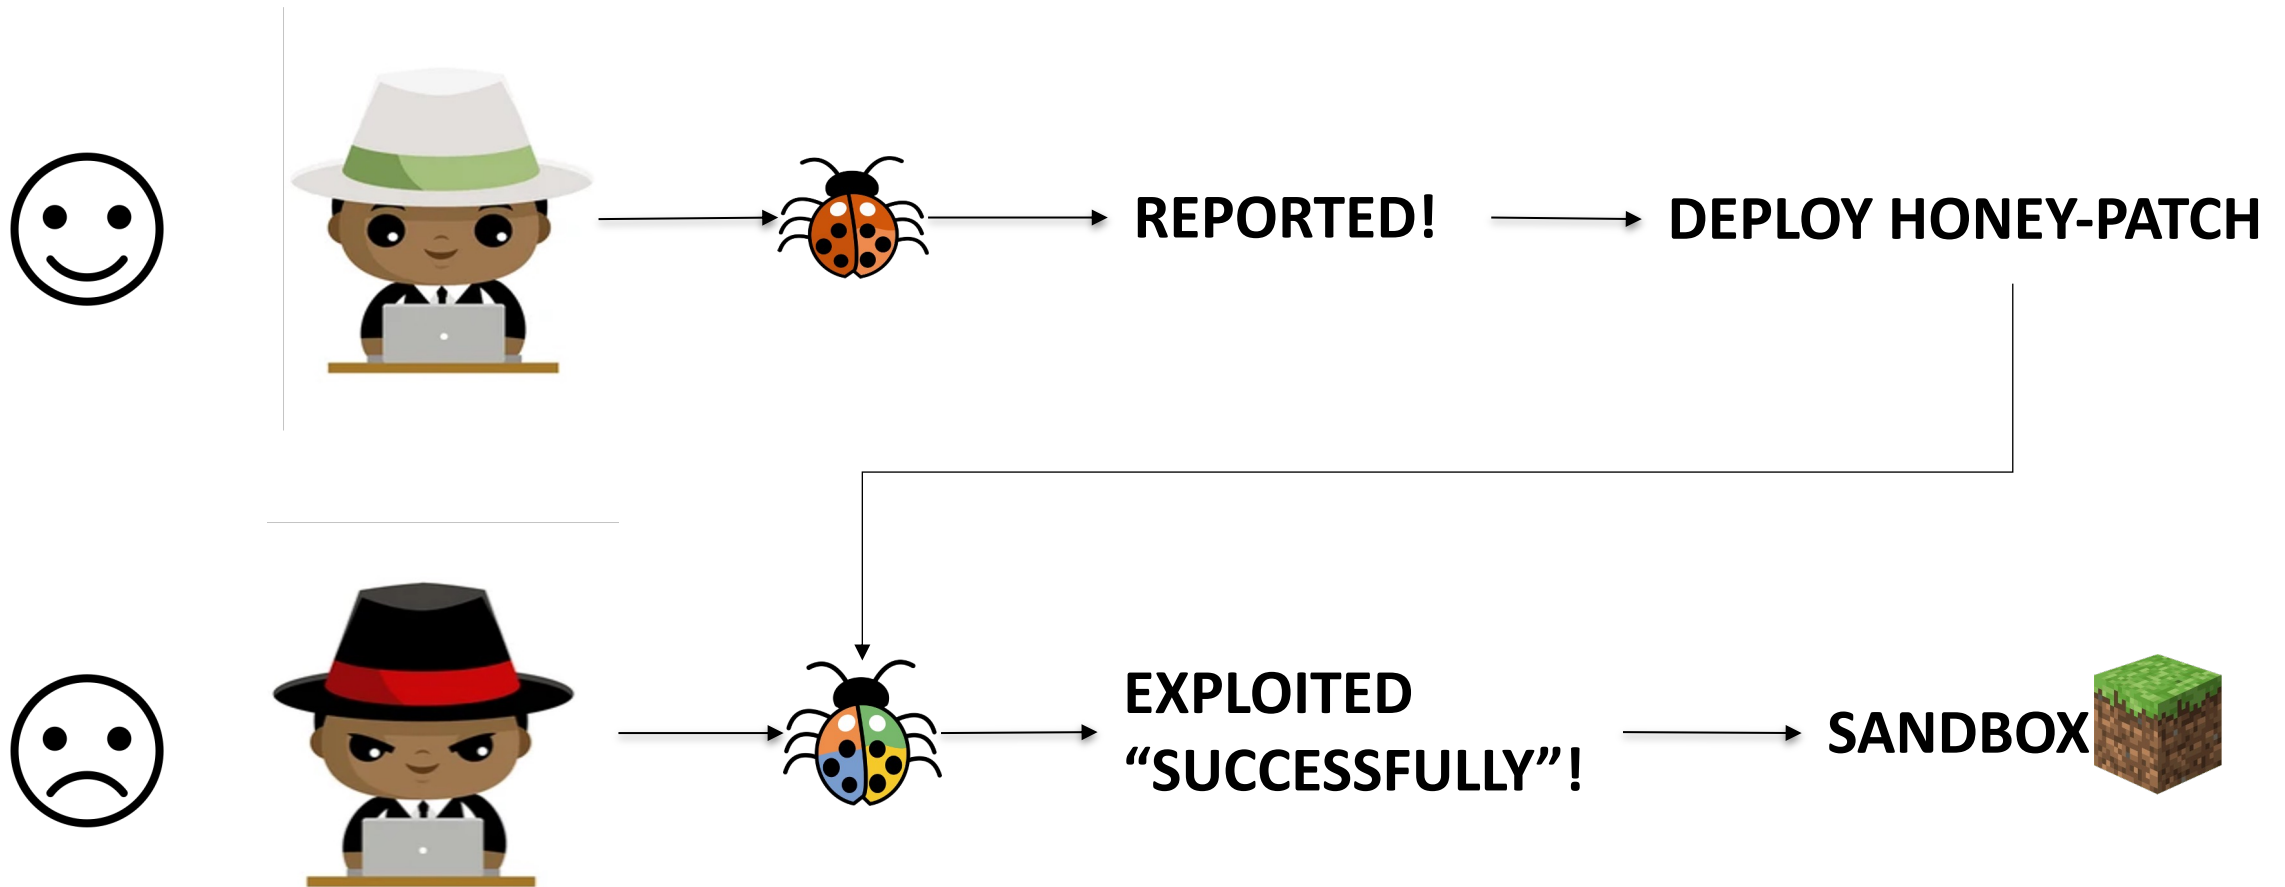
REPORTED!

**DEPLOY
UPDATES/PATCHES**



EXPLOITED!

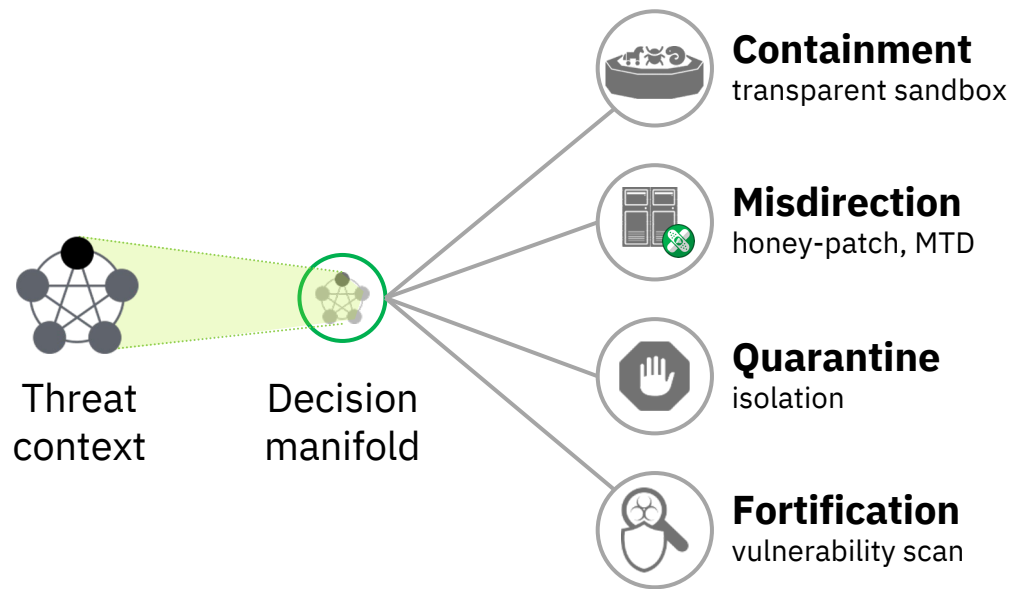
Using Deception as Defense



Agenda

- ❖ Reviewing Honey-Patching
- ❖ Setting up the user-study
- ❖ Deploying Honey-Patches as mitigations
- ❖ Modeling the adversary's attack graph
- ❖ Formulating the Markov Game model
- ❖ Computing the Bayesian Stackelberg Equilibrium

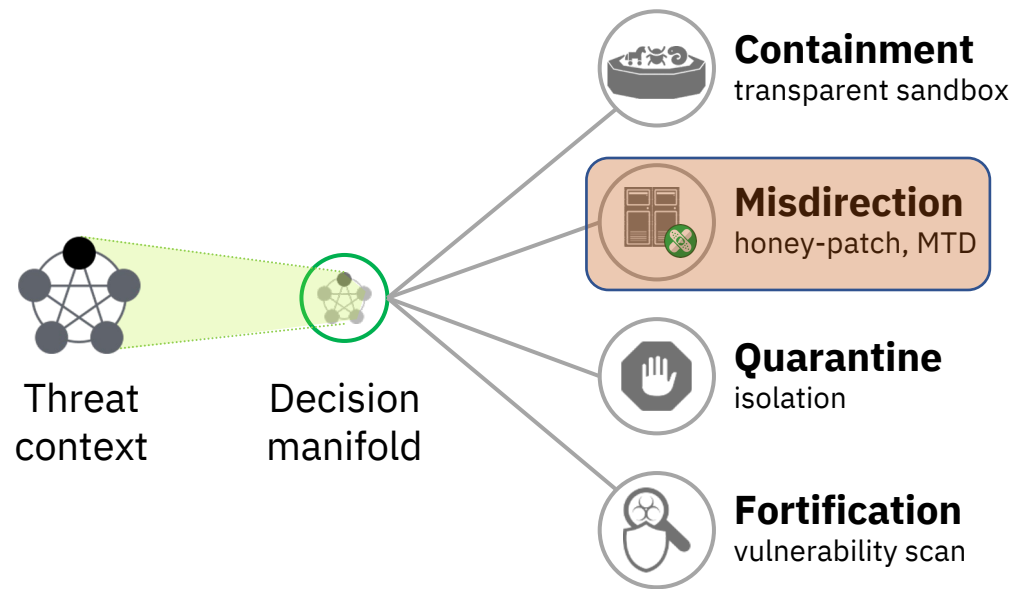
Automating Proactive Responses



Which one
to choose?

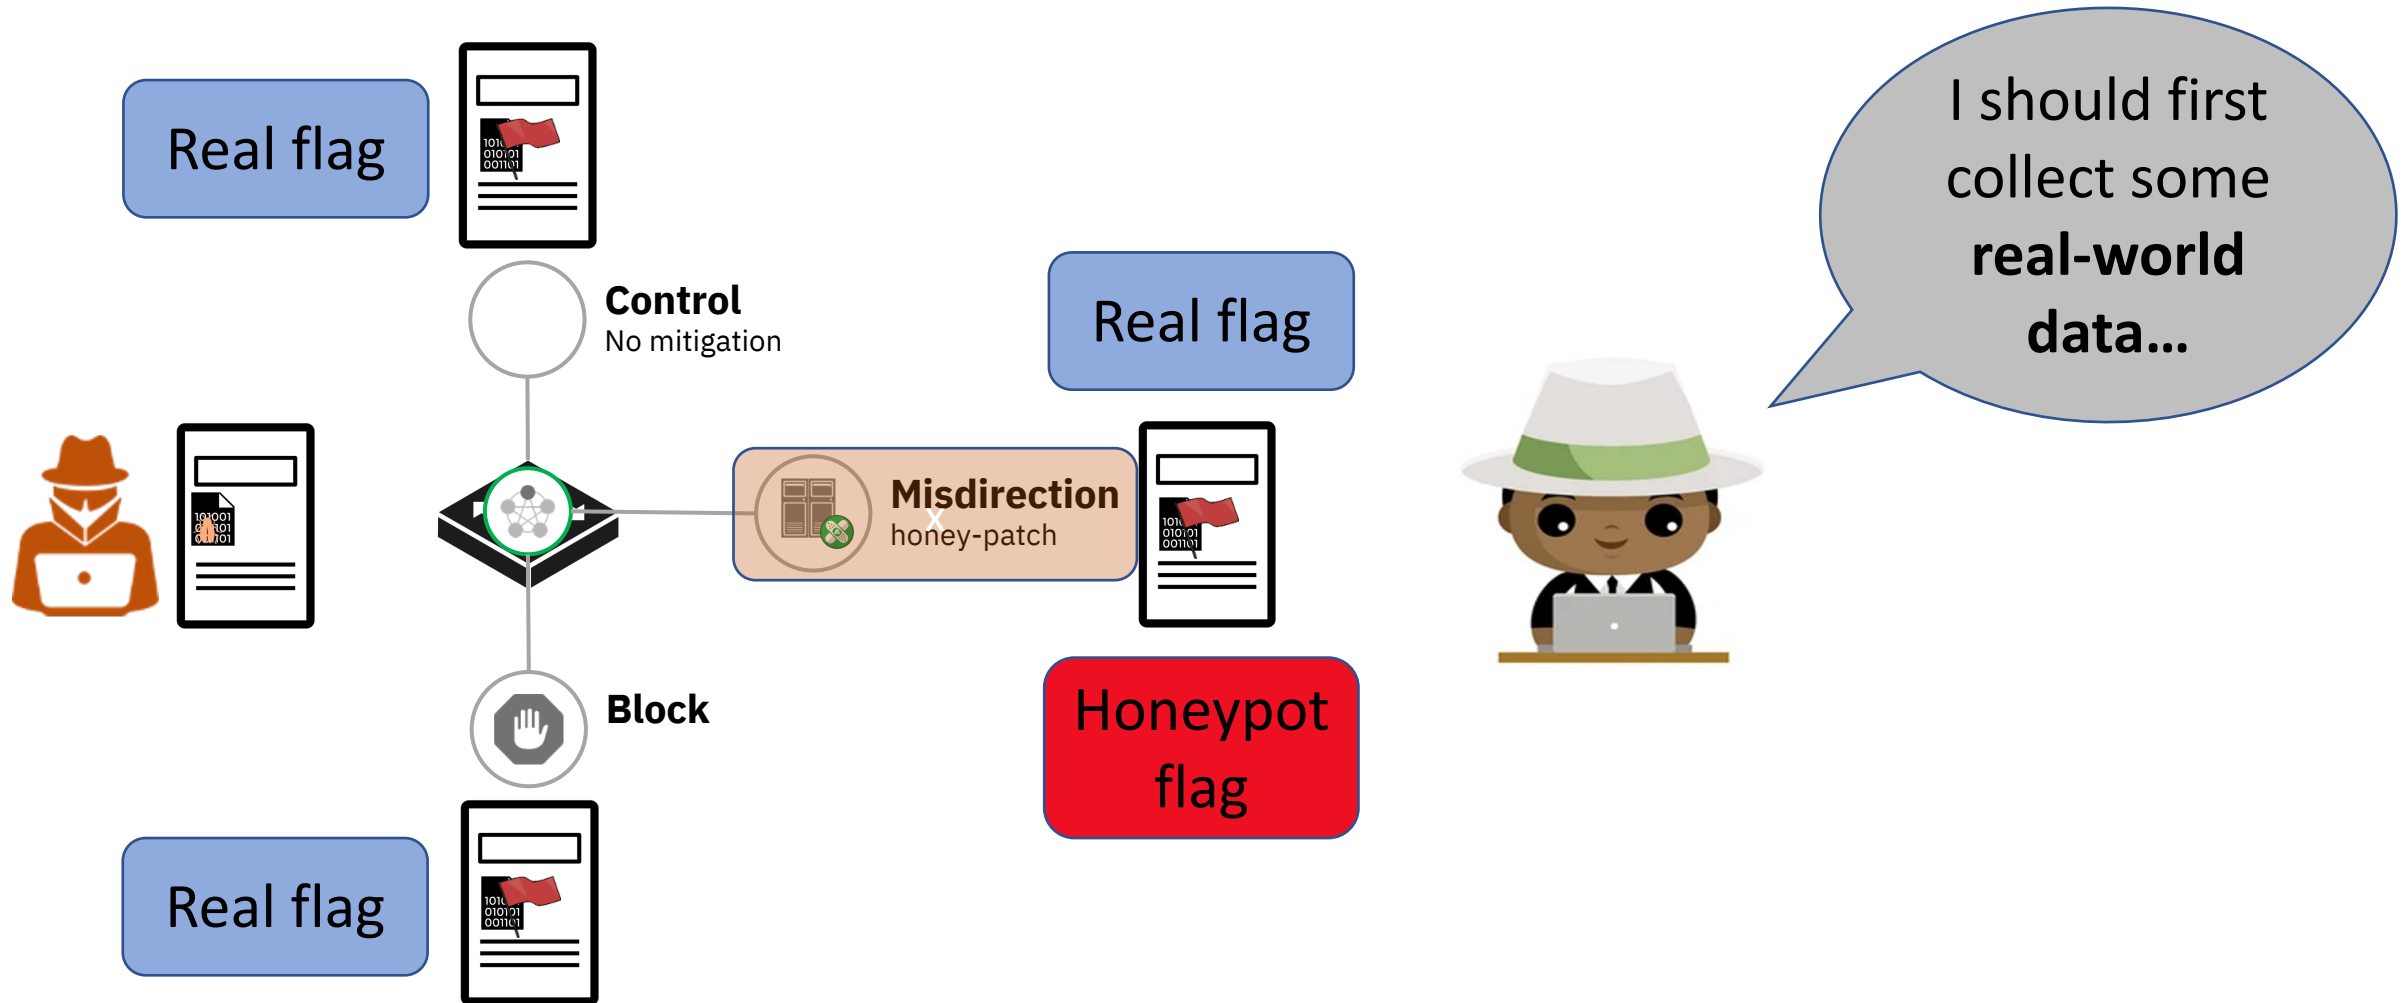
I want to **disrupt**
attack kill chains
and perform
attacker
reconnaissance!

Automating Proactive Responses

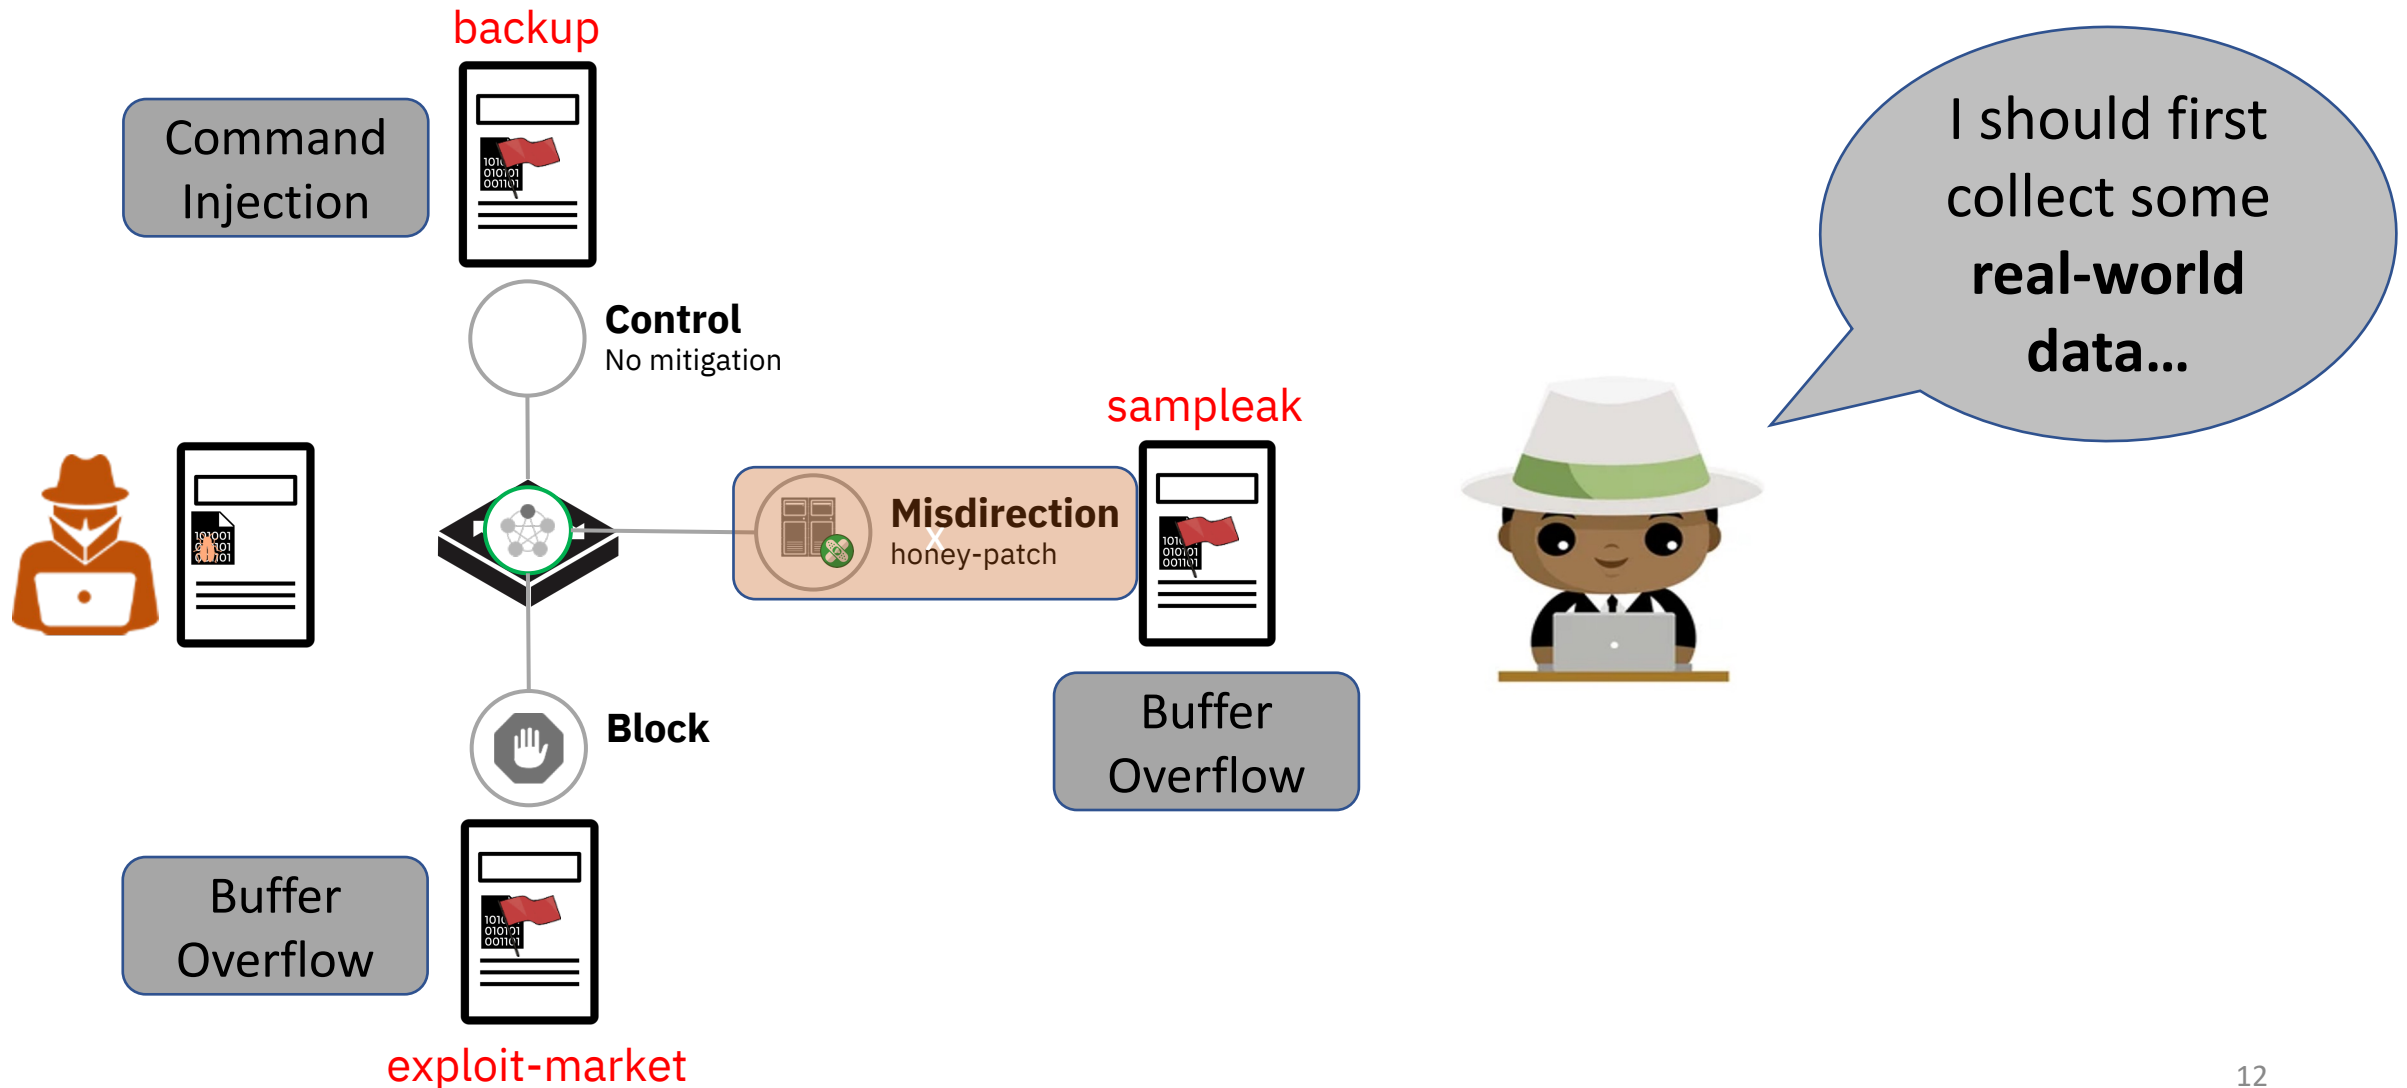


I also need to
understand the
attacker's
behavior!

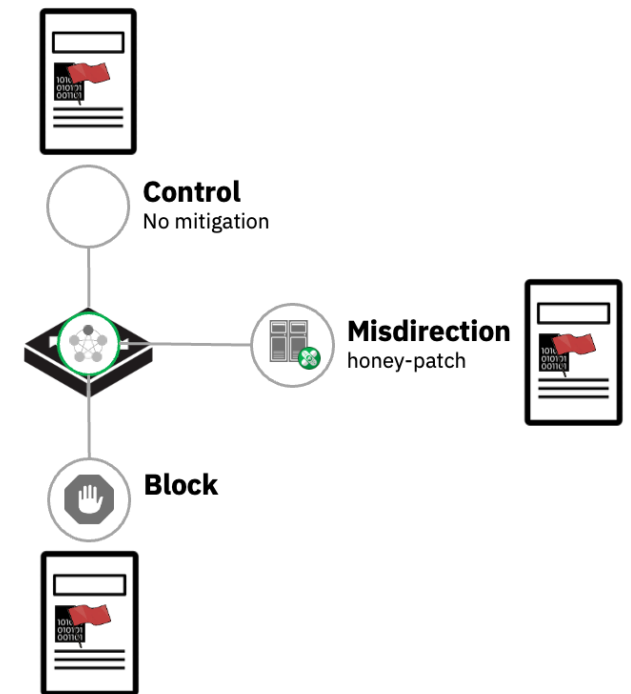


Testing Deception as Defense Strategy



Testing Deception as Defense Strategy

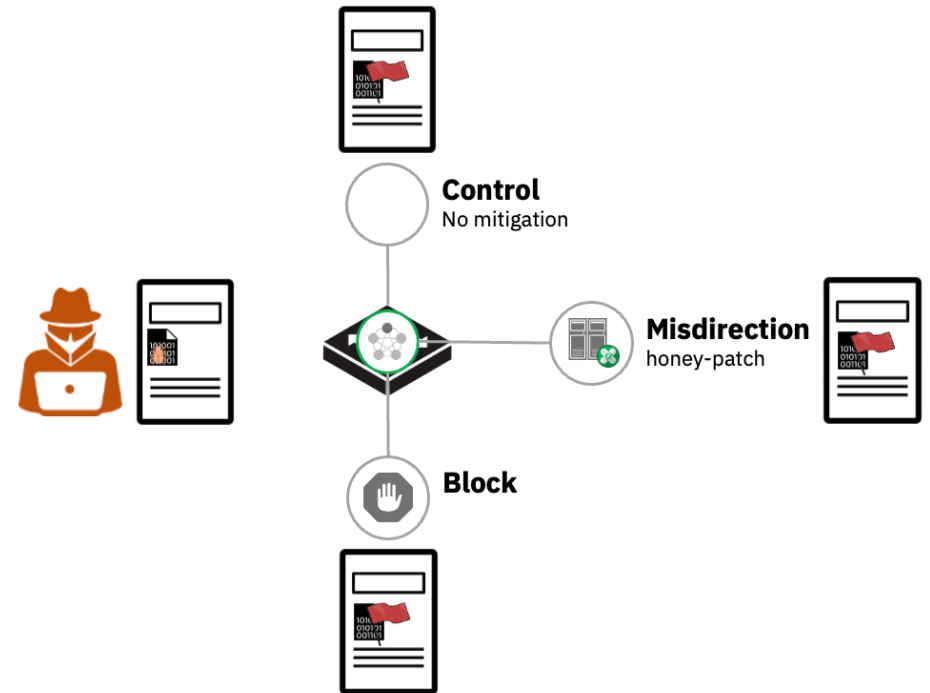


Adapting the iCTF Framework



Hypothesis Testing

“Once trapped in a honeypot environment, the attacker chooses to continue with the existing strategy to exploit the remaining vulnerabilities.”




Agenda

- ❖ Reviewing Honey-Patching
- ❖ Setting up the user-study
- ❖ Deploying Honey-Patches as mitigations
- ❖ Modeling the adversary's attack graph
- ❖ Formulating the Markov Game model
- ❖ Computing the Bayesian Stackelberg Equilibrium


Deploying the Honey-Patching Framework

1. Patch synthesis
2. Inject patch and compile




Deploying the Honey-Patching Framework

1. Patch synthesis
2. Inject patch and compile
3. Setup trampolines



Deploying the Honey-Patching Framework

1. Patch synthesis
2. Inject patch and compile
3. Setup trampolines
4. Booby trap triggered
5. Sandbox attacker





INSIDER Honey-Patching Framework

Deploying the Honey-Patching Framework

Original Chall



```
FROM ubuntu:18.04
...
CMD cd service/rw && socat tcp-l:6666,reuseaddr,fork
exec:"../ro/sampleak"
```



Deploying the Honey-Patching Framework

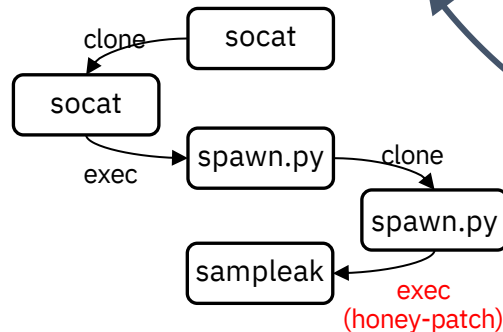
Original Chall

```
FROM ubuntu:18.04
...
CMD cd service/rw && socat tcp-l:6666,reuseaddr,fork
exec:"../ro/sampleak"
```



Honey-Patched Chall

```
FROM ubuntu:18.04
...
CMD cd service/rw && socat tcp-l:6666,reuseaddr,fork
exec:"/insider/spawn.py patch.conf ../ro/sampleak"
```

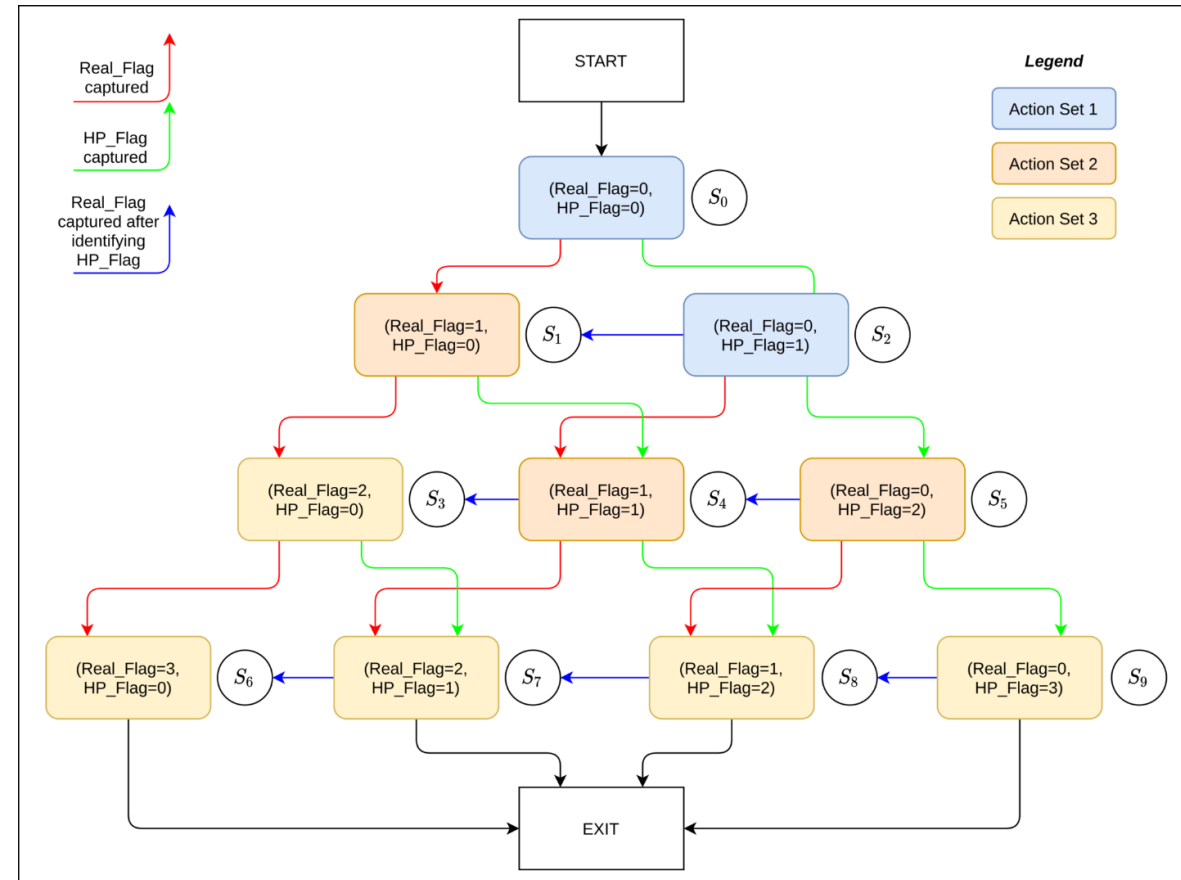
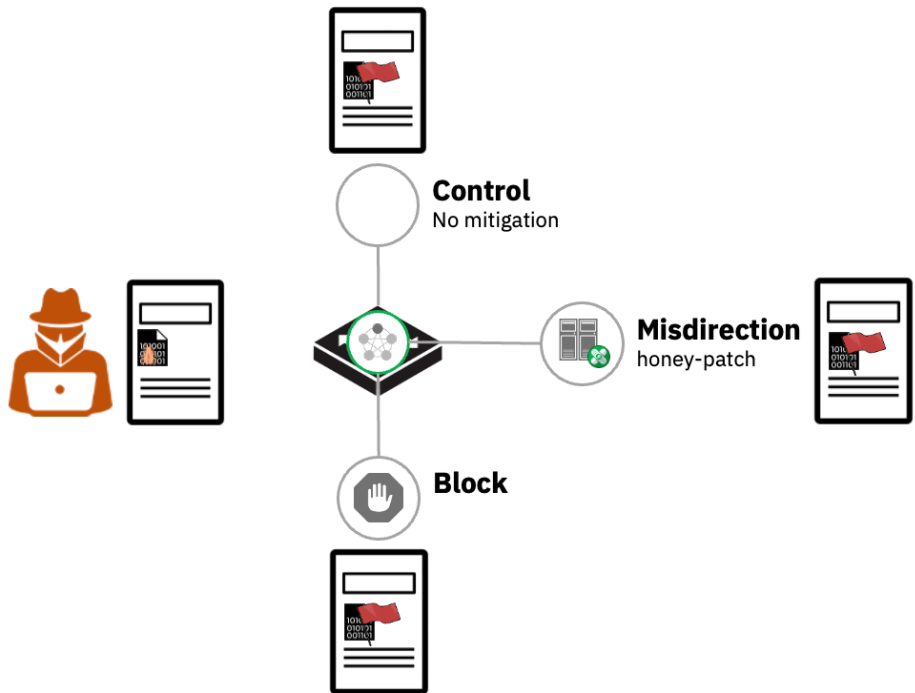


```
{
  "__type__": "config",
  "agentscript": "patch.js",
  "spawnscript": "spawn.js",
  "elfpaths": [
    {
      "path": "/insider/frida/examples/sampleak/sampleak",
      "pic": false
    }
  ],
  "bcpath": "patch.ll",
  "replacements": [
    {
      "targetfn": "is_easy_password",
      "surrogatefn": "_p_is_easy_password",
      "fntype": { "rettype": "int", "argstype": ["pointer"] }
    }
  ],
  "logpath": "/insider/frida/examples/sampleak/spawn_logs",
  "target": "sampleak"
}
```

Agenda

- ❖ Reviewing Honey-Patching
- ❖ Setting up the user-study
- ❖ Deploying Honey-Patches as mitigations
- ❖ Modeling the adversary's attack graph
- ❖ Formulating the Markov Game model
- ❖ Computing the Bayesian Stackelberg Equilibrium

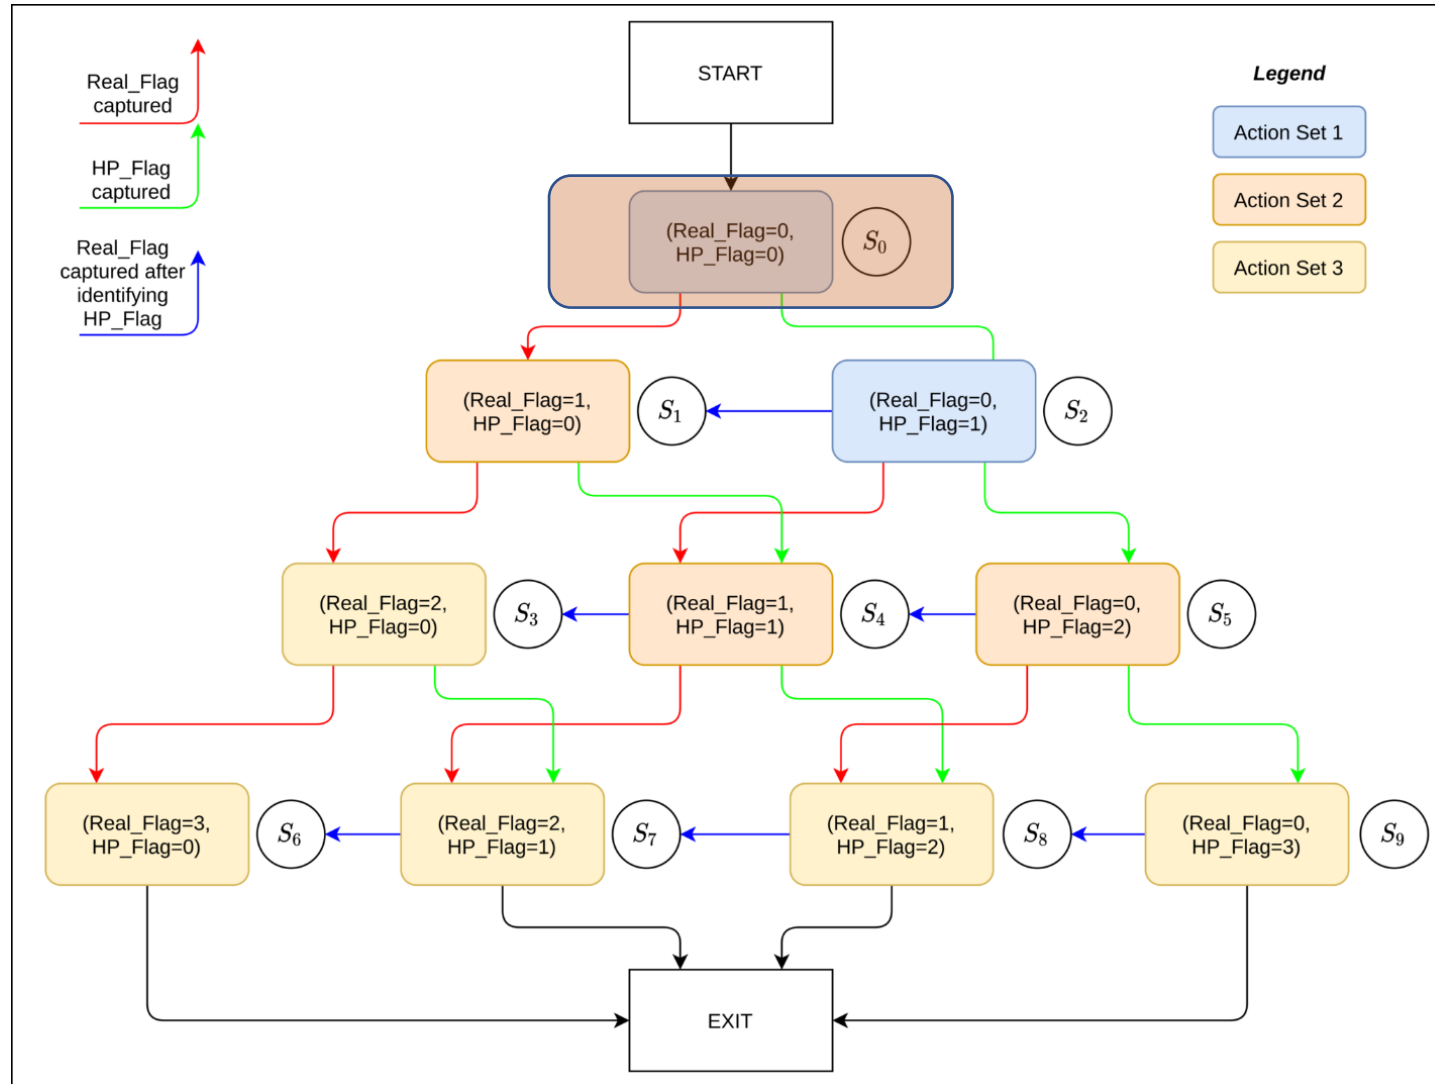
Attack Graph



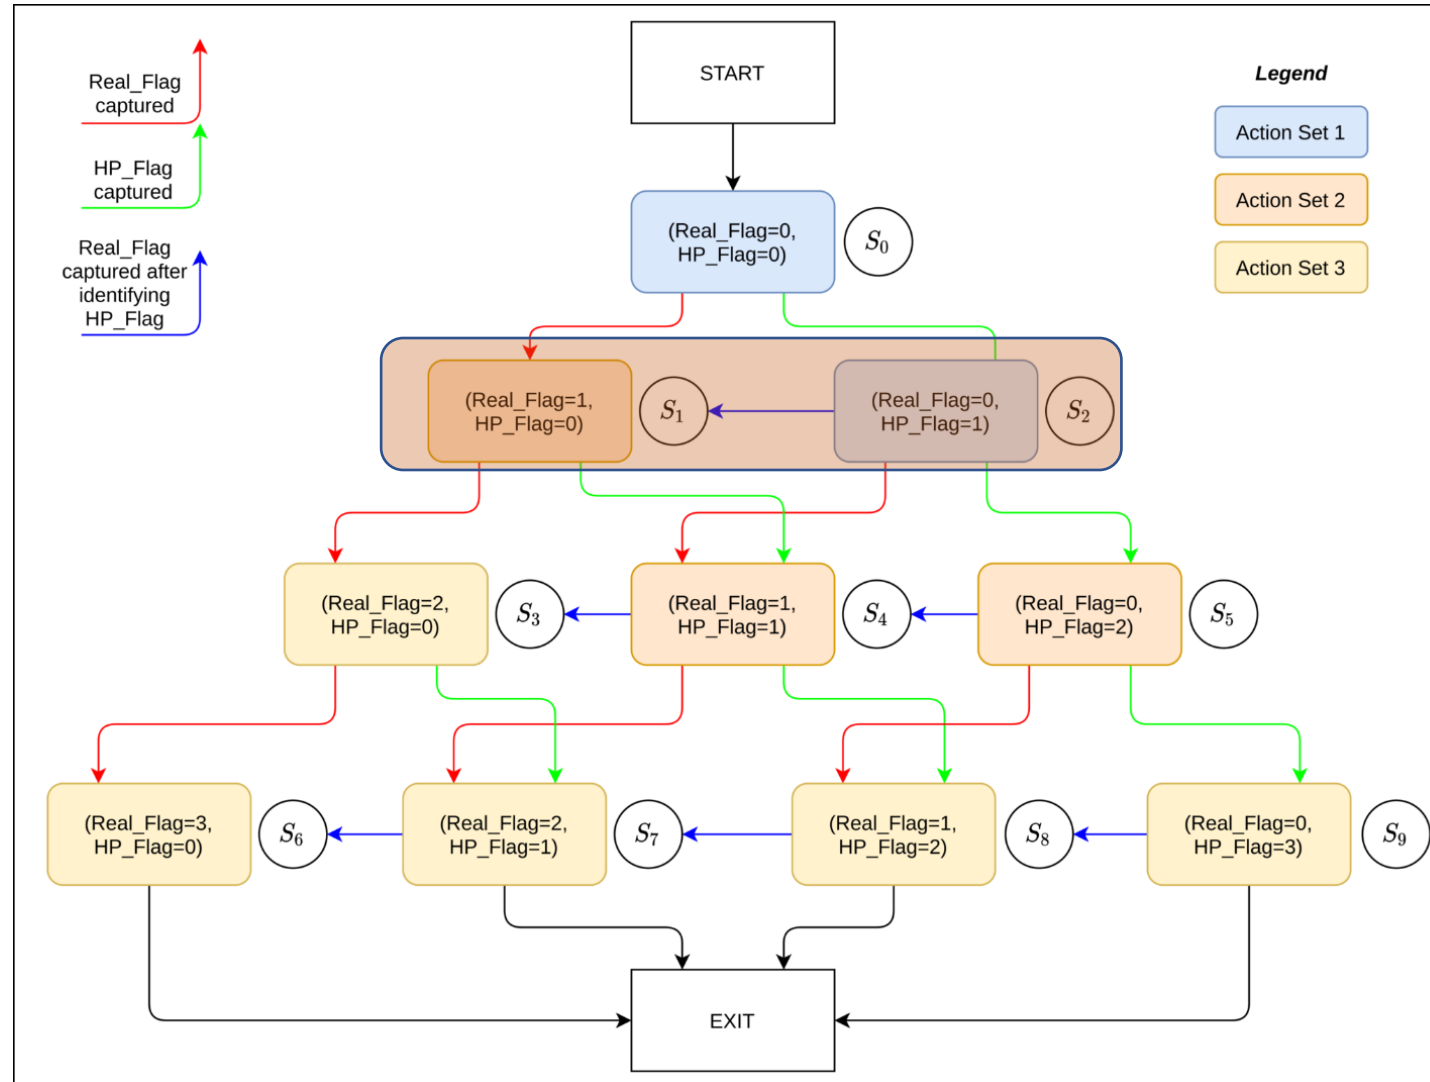
Attack Graph

Let me
start the
exploit!

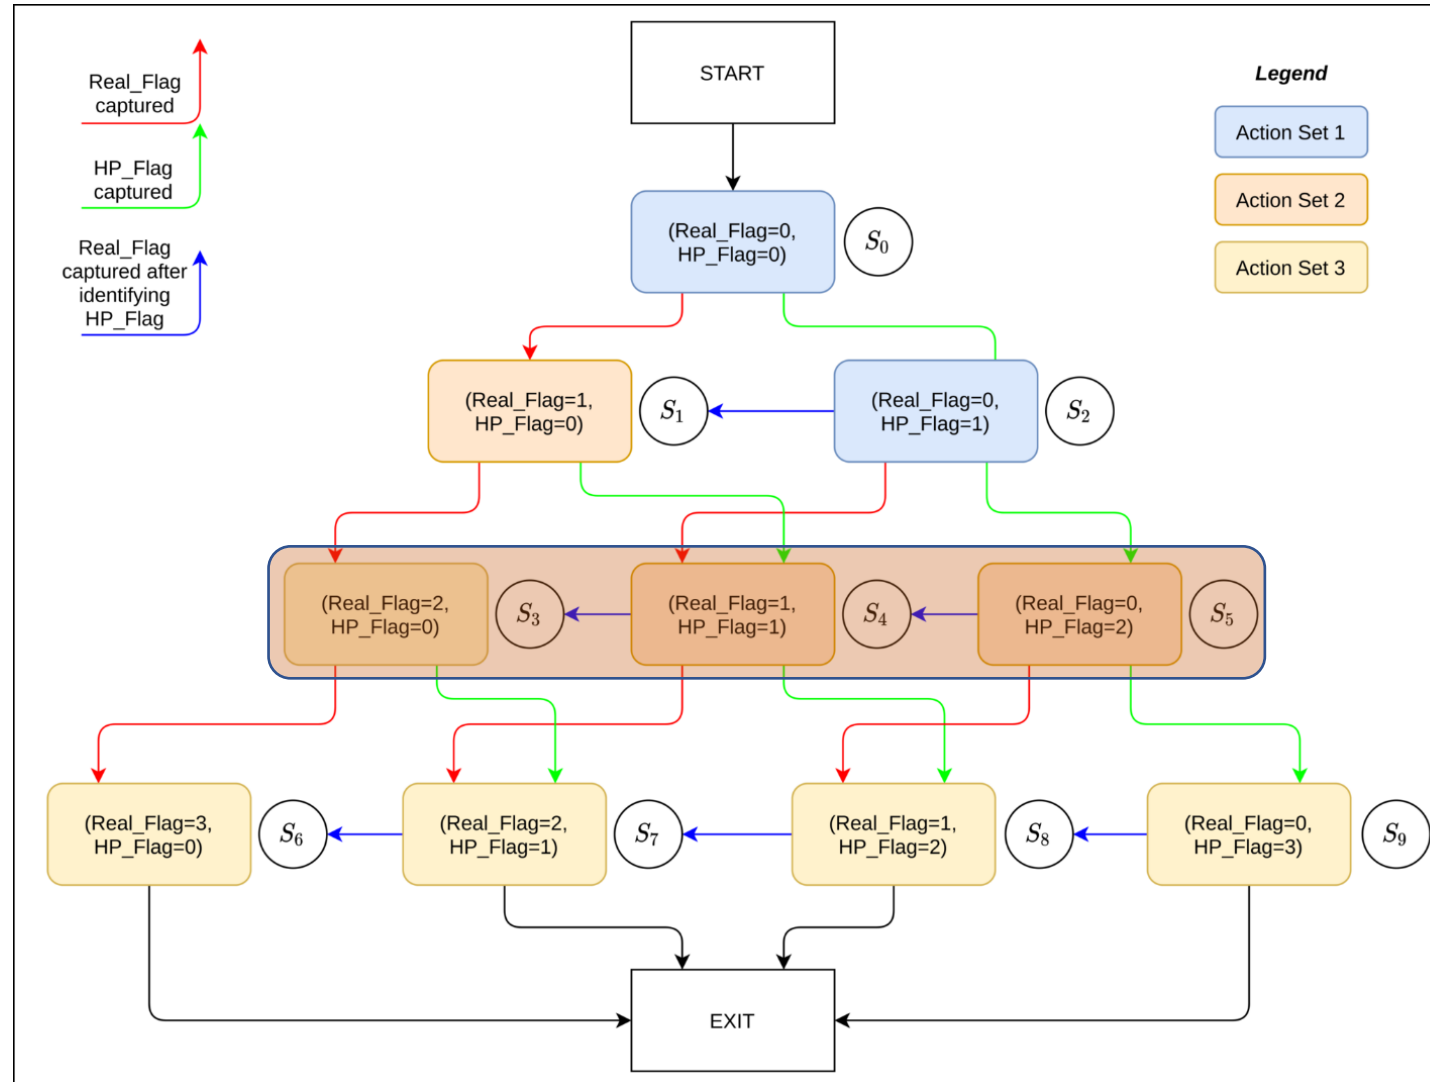
Which
flag is it
going to
be?



Attack Graph



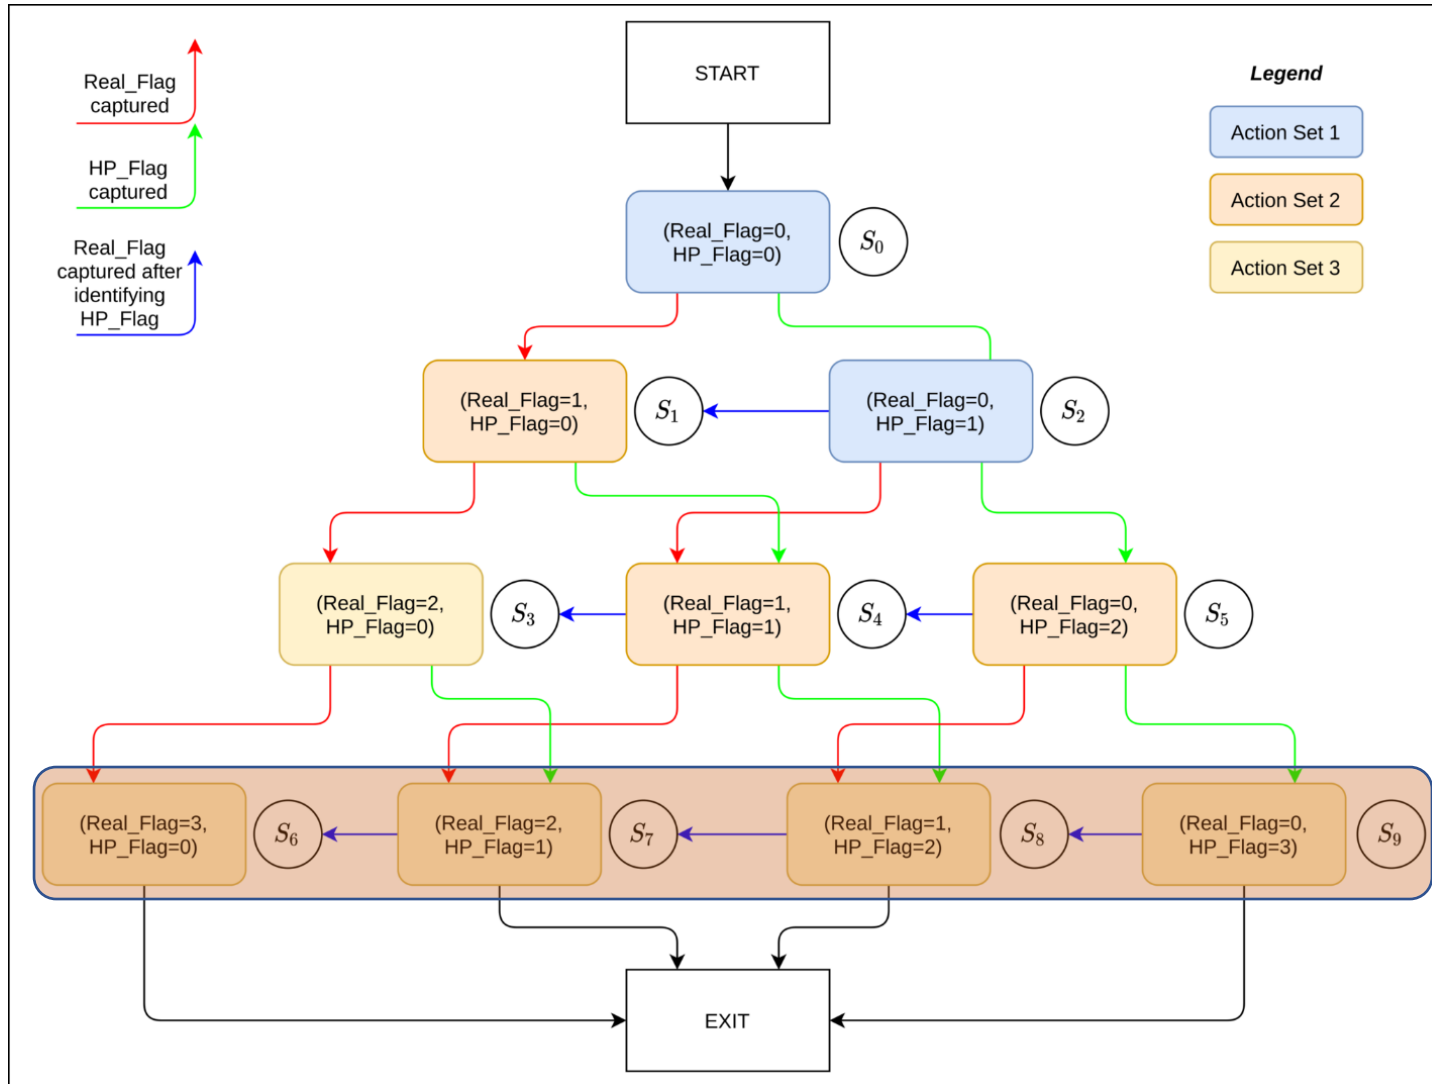
Attack Graph



Attack Graph

I got the
real
flags!

Are you
sure?




Agenda

- ❖ Reviewing Honey-Patching
- ❖ Setting up the user-study
- ❖ Deploying Honey-Patches as mitigations
- ❖ Modeling the adversary's attack graph
- ❖ Formulating the Markov Game model
- ❖ Computing the Bayesian Stackelberg Equilibrium

Formulating the Markov Game

Markov Game (Shapley 1953) for two players P_1 and P_2 can be defined by the tuple $(S, A_1, A_2, \tau, R, \gamma)$ where,


- $S = \{s_1, s_2, \dots, s_k\}$ represents a set of finite states of the game,



Formulating the Markov Game

Markov Game (Shapley 1953) for two players P_1 and P_2 can be defined by the tuple $(S, A_1, A_2, \tau, R, \gamma)$ where,


- $S = \{s_1, s_2, \dots, s_k\}$ represents a set of finite states of the game,
- $A_1 = \{a^1_1, a^2_1, \dots, a^m_1\}$ represents the possible finite action sets for P_1 ,
- $A_2 = \{a^1_2, a^2_2, \dots, a^m_2\}$ represents the possible finite action sets for P_2 ,



Formulating the Markov Game

Markov Game (Shapley 1953) for two players P_1 and P_2 can be defined by the tuple $(S, A_1, A_2, \tau, R, \gamma)$ where,


- $S = \{s_1, s_2, \dots, s_k\}$ represents a set of finite states of the game,
- $A_1 = \{a^1_1, a^2_1, \dots, a^m_1\}$ represents the possible finite action sets for P_1 ,
- $A_2 = \{a^1_2, a^2_2, \dots, a^m_2\}$ represents the possible finite action sets for P_2 ,
- $\tau(s, a_1, a_2, s')$ represents the transition probability matrix from a state $s \in S$ to $s' \in S$ when the P_1 and P_2 take actions a_1 and a_2 respectively,




Formulating the Markov Game

Markov Game (Shapley 1953) for two players P_1 and P_2 can be defined by the tuple $(S, A_1, A_2, \tau, R, \gamma)$ where,


- $S = \{s_1, s_2, \dots, s_k\}$ represents a set of finite states of the game,
- $A_1 = \{a^1_1, a^2_1, \dots, a^m_1\}$ represents the possible finite action sets for P_1 ,
- $A_2 = \{a^1_2, a^2_2, \dots, a^m_2\}$ represents the possible finite action sets for P_2 ,
- $\tau(s, a_1, a_2, s')$ represents the transition probability matrix from a state $s \in S$ to $s' \in S$ when the P_1 and P_2 take actions a_1 and a_2 respectively,
- $R^i(s, a_1, a_2)$ denotes the utility or the rewards received by P_i in state s when P_1 and P_2 take actions a_1 and a_2 respectively,
- $\gamma \rightarrow [0, 1)$ is discount factor for future rewards.



Formulating the Markov Game




Attacker's actions (A_A)



Defender's actions (A_D)

	no_mon	hp_1	hp_2	hp_3
no_op	0	-3	-3	-3
exp_1	-5.9	2.9	-8.9	-8.9
exp_2	-5.9	-8.9	2.9	-8.9
exp_3	-5.9	-8.9	-8.9	2.9

Deploy Misdirection



Game Theory Model

Game Theoretic Modeling – Updating the Parameters

1. Naive Model

- Keeping all other parameters constant, we **randomly set the transition probabilities τ** for the game-theoretic model.

Game Theoretic Modeling – Updating the Parameters

1. Naive Model

- Keeping all other parameters constant, we **randomly set the transition probabilities τ** for the game-theoretic model.

2. Naive Model

- Keeping all other parameters constant, **a system expert sets the transition probabilities τ** for the game-theoretic model.

Game Theoretic Modeling – Updating the Parameters

1. Naive Model

- Keeping all other parameters constant, we **randomly set the transition probabilities τ** for the game-theoretic model.

2. Naive Model

- Keeping all other parameters constant, **a system expert sets the transition probabilities τ** for the game-theoretic model.


3. Updated Model

- Transition probability matrix τ **is calculated using the statistics** obtained from the iCTF user studies.

Game Theoretic Modeling – Updating the Parameters

1. Naive Model

- Keeping all other parameters constant, we **randomly set the transition probabilities τ** for the game-theoretic model.




2. Naive Model

- Keeping all other parameters constant, **a system expert sets the transition probabilities τ** for the game-theoretic model.



3. Updated Model

- Transition probability matrix **τ is calculated using the statistics** obtained from the iCTF user studies.



τ

Agenda

- ❖ Reviewing Honey-Patching
- ❖ Setting up the user-study
- ❖ Deploying Honey-Patches as mitigations
- ❖ Modeling the adversary's attack graph
- ❖ Formulating the Markov Game model
- ❖ Computing the Bayesian Stackelberg Equilibrium

Finding the Optimal Defender Strategy

OPT
does the
best!



Algorithms:

- **MMP**: Min Max Pure Strategy
- **URS**: Uniform Random Strategy
- **OPT**: Optimal Mixed Strategy






Figure: Defender's payoffs for Naive Model - **randomly set transition probabilities** and uniform mitigation deployment costs.




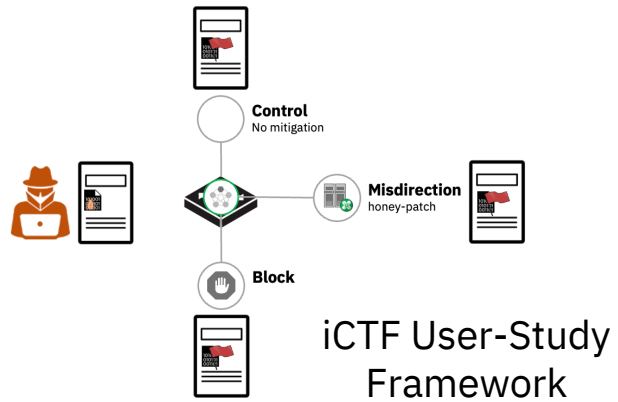



Figure: Defender's payoffs for Naive Model – **system expert set transition probabilities** and uniform mitigation deployment costs.

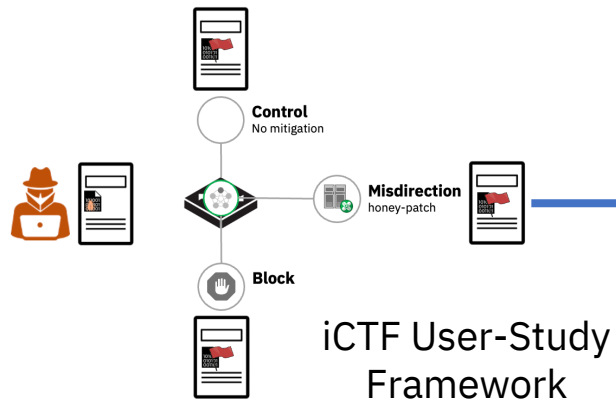
Updating the Model using iCTF User Studies



Updating the Model using iCTF User Studies



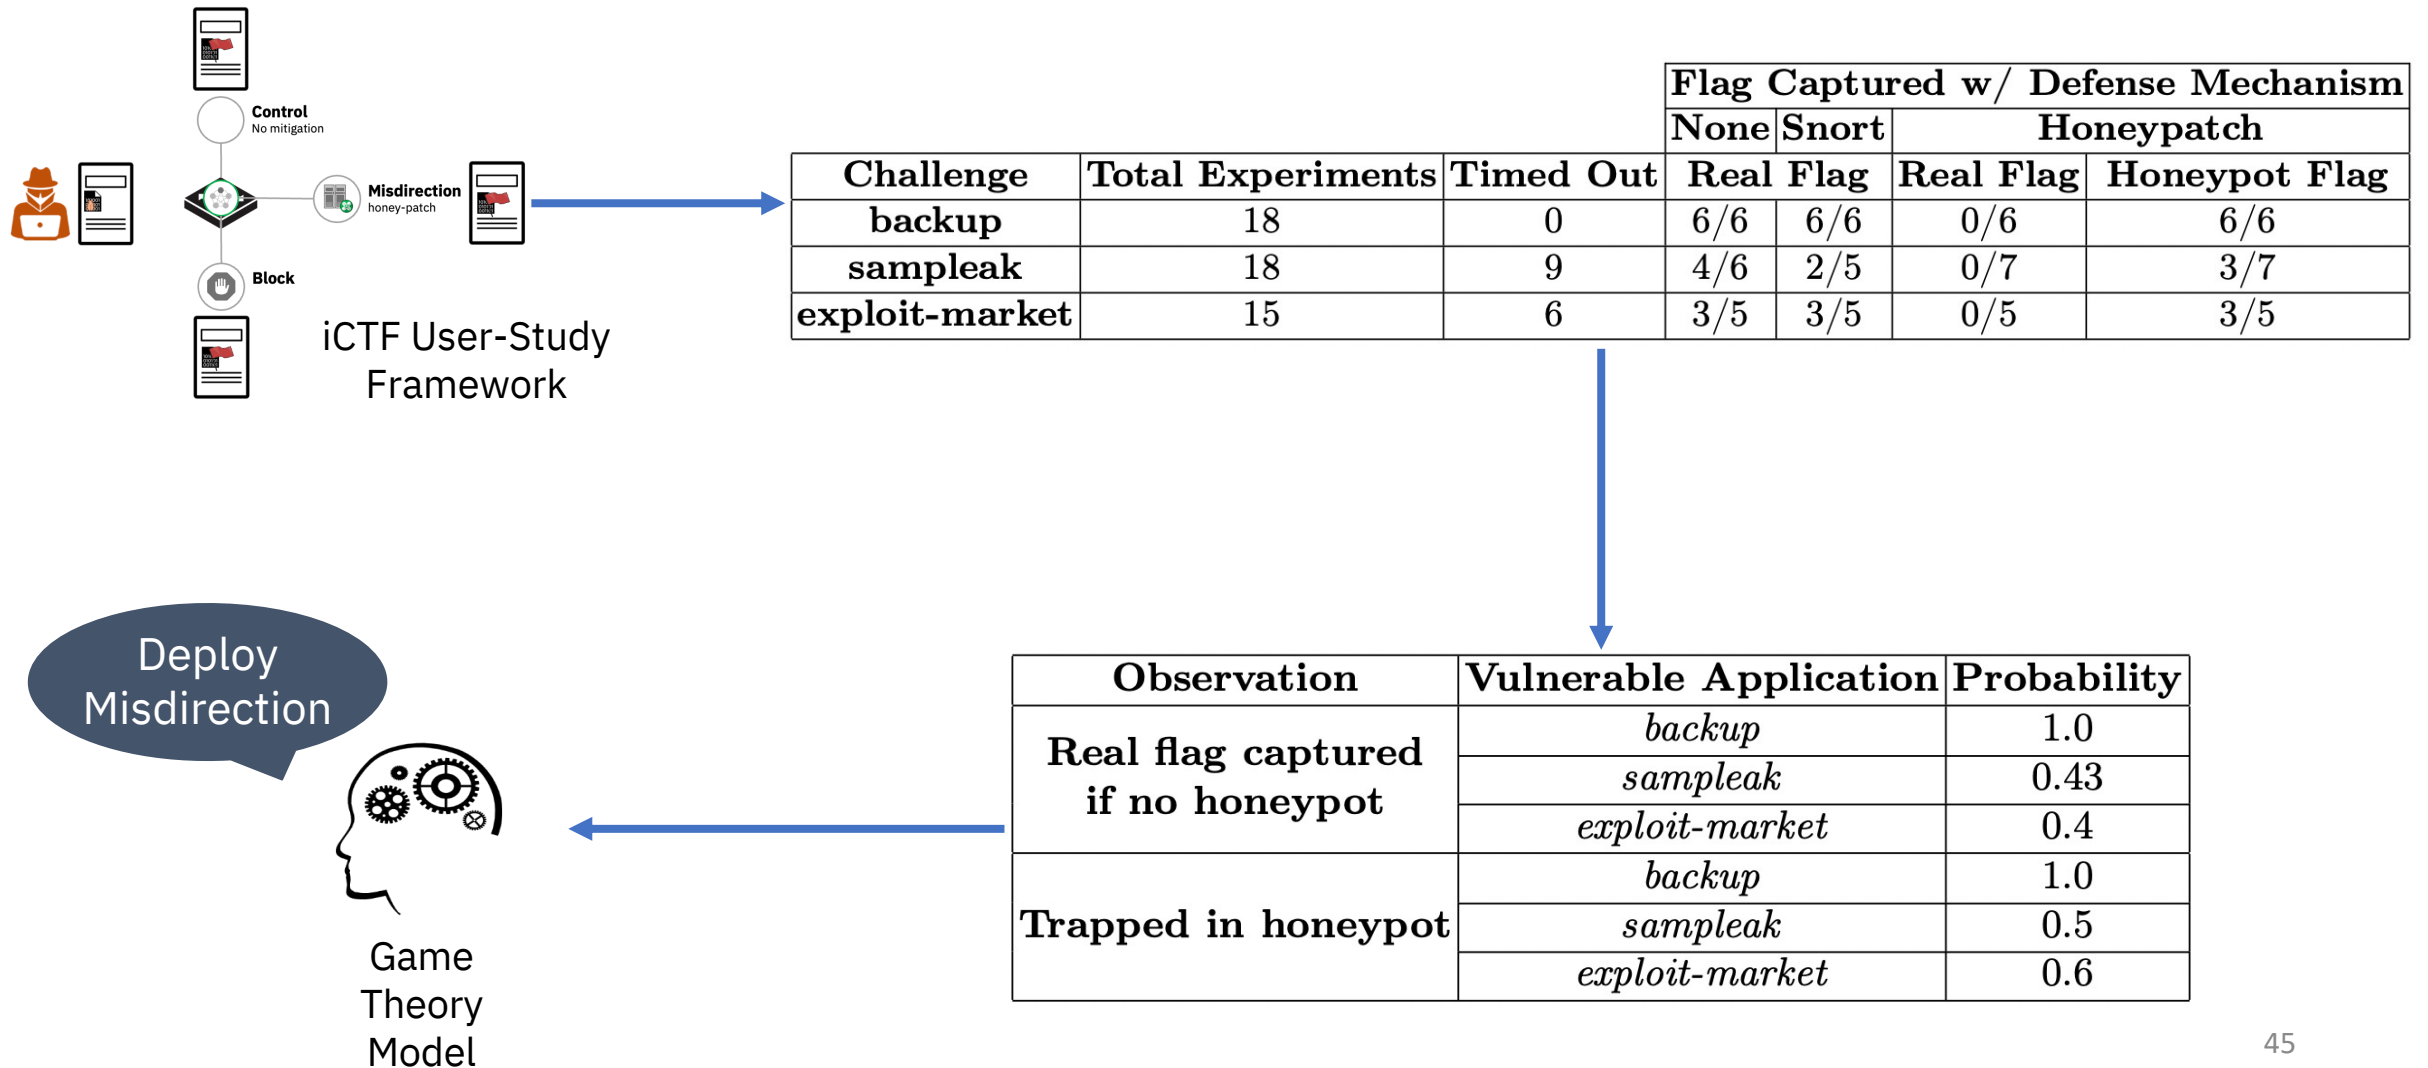
Updating the Model using iCTF User Studies



Challenge	Total Experiments	Timed Out	Flag Captured w/ Defense Mechanism			
			None	Snort	Honeypatch	
			Real Flag		Real Flag	Honeypot Flag
backup	18	0	6/6	6/6	0/6	6/6
sampleak	18	9	4/6	2/5	0/7	3/7
exploit-market	15	6	3/5	3/5	0/5	3/5

Observation	Vulnerable Application	Probability
Real flag captured if no honeypot	<i>backup</i>	1.0
	<i>sampleak</i>	0.43
	<i>exploit-market</i>	0.4
Trapped in honeypot	<i>backup</i>	1.0
	<i>sampleak</i>	0.5
	<i>exploit-market</i>	0.6

Updating the Model using iCTF User Studies



Comparing the Game Theoretic Model Case Studies




Figure: Defender's payoffs compared for the three models using **uniform mitigation deployment costs**.




Figure: Defender's payoffs compared for the three models using **non-uniform mitigation deployment costs**.

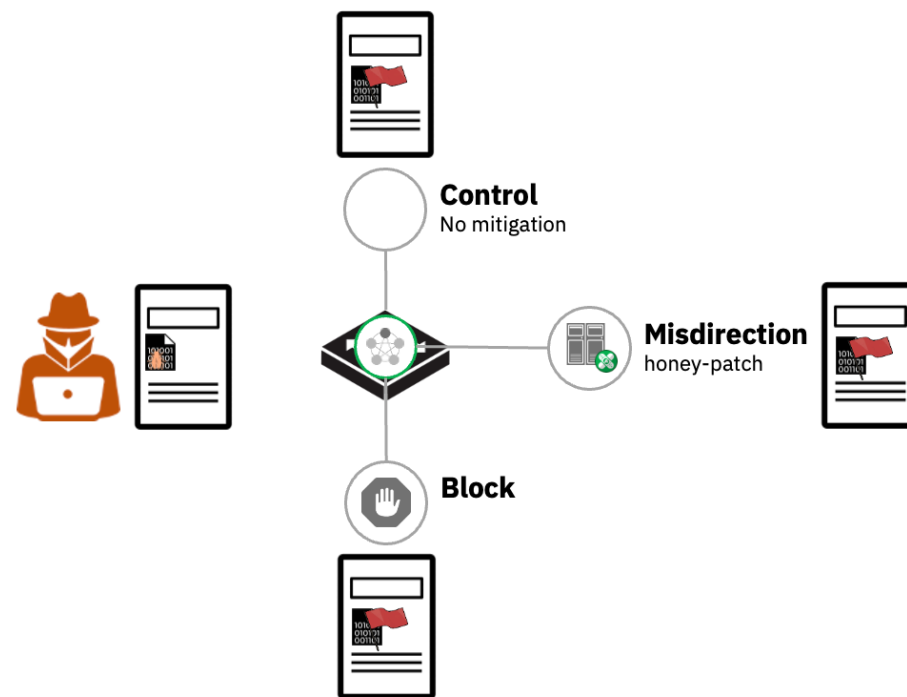


Defender's returns:

Expert set (τ) > Randomly set (τ) > Computed from iCTF user studies (τ)

Evaluating the Hypothesis

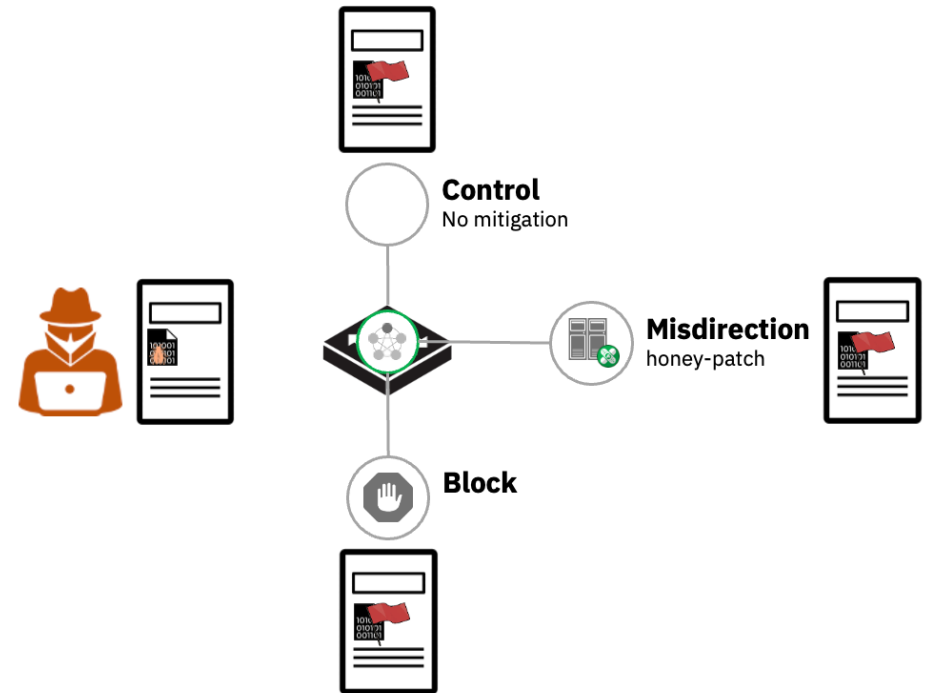
Hypothesis: Once trapped in a honeypot environment, the attacker chooses to continue with the existing strategy to exploit the remaining vulnerabilities.



Evaluating the Hypothesis

Hypothesis: Once trapped in a honeypot environment, the attacker chooses to continue with the existing strategy to exploit the remaining vulnerabilities.

Observation 1: None of the attackers received the observation of being trapped in a honey-pot, and thus continued with their existing strategy.

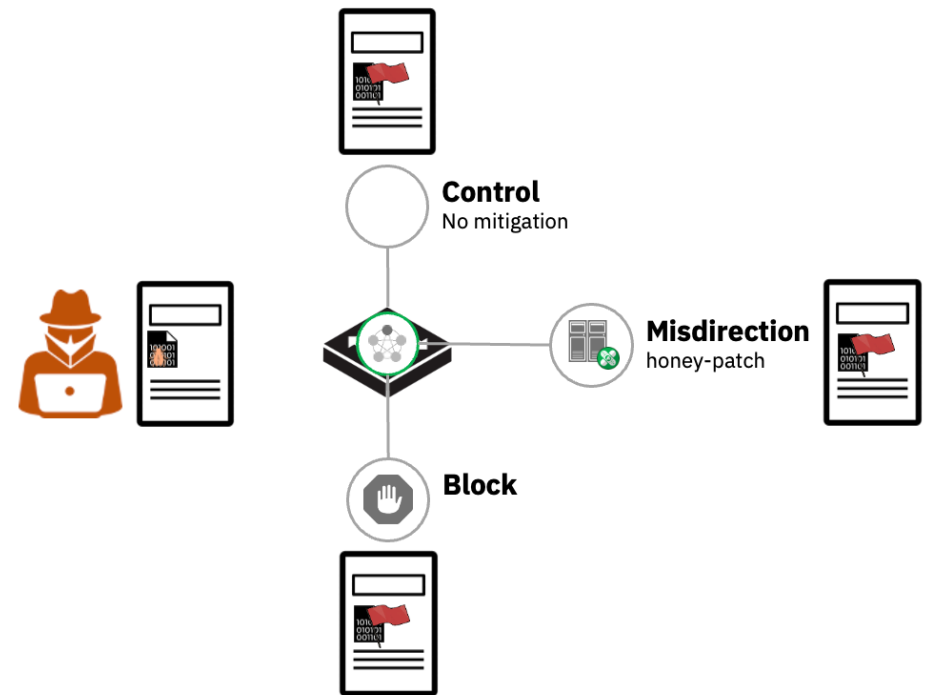


Evaluating the Hypothesis

Hypothesis: Once trapped in a honeypot environment, the attacker chooses to continue with the existing strategy to exploit the remaining vulnerabilities.

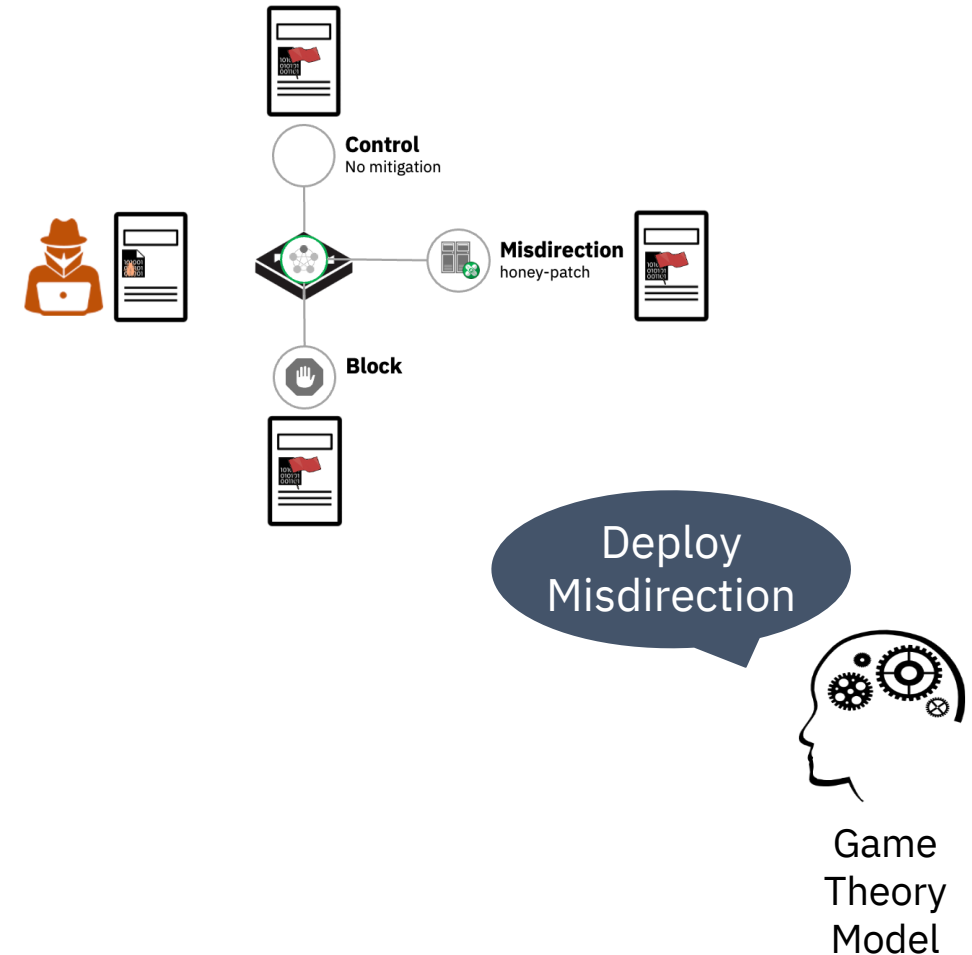
Observation 1: None of the attackers received the observation of being trapped in a honey-pot, and thus continued with their existing strategy.

Observation 2: Only in one instance, an attacker learns about the honeypot for sample leak, after the user study ends and the attacker is explicitly informed.



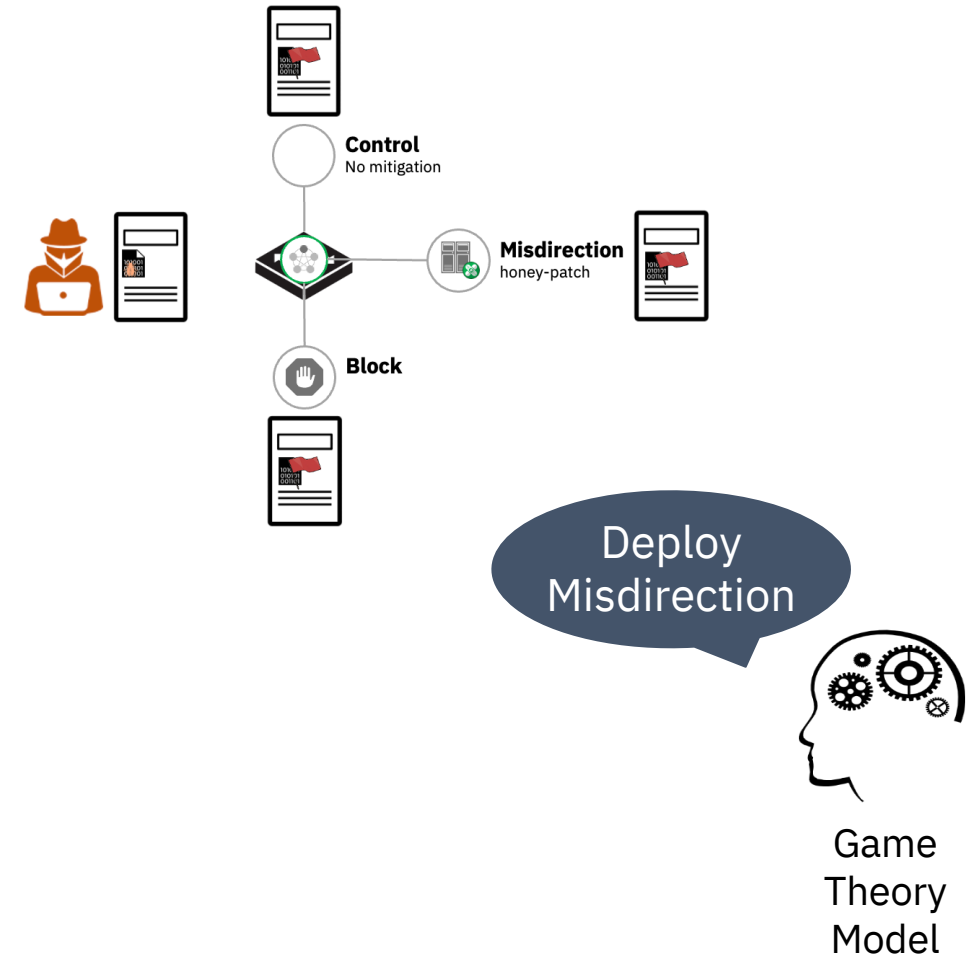
Conclusion

- Utility returns vary in the three models only for the initial 3 states of the game.



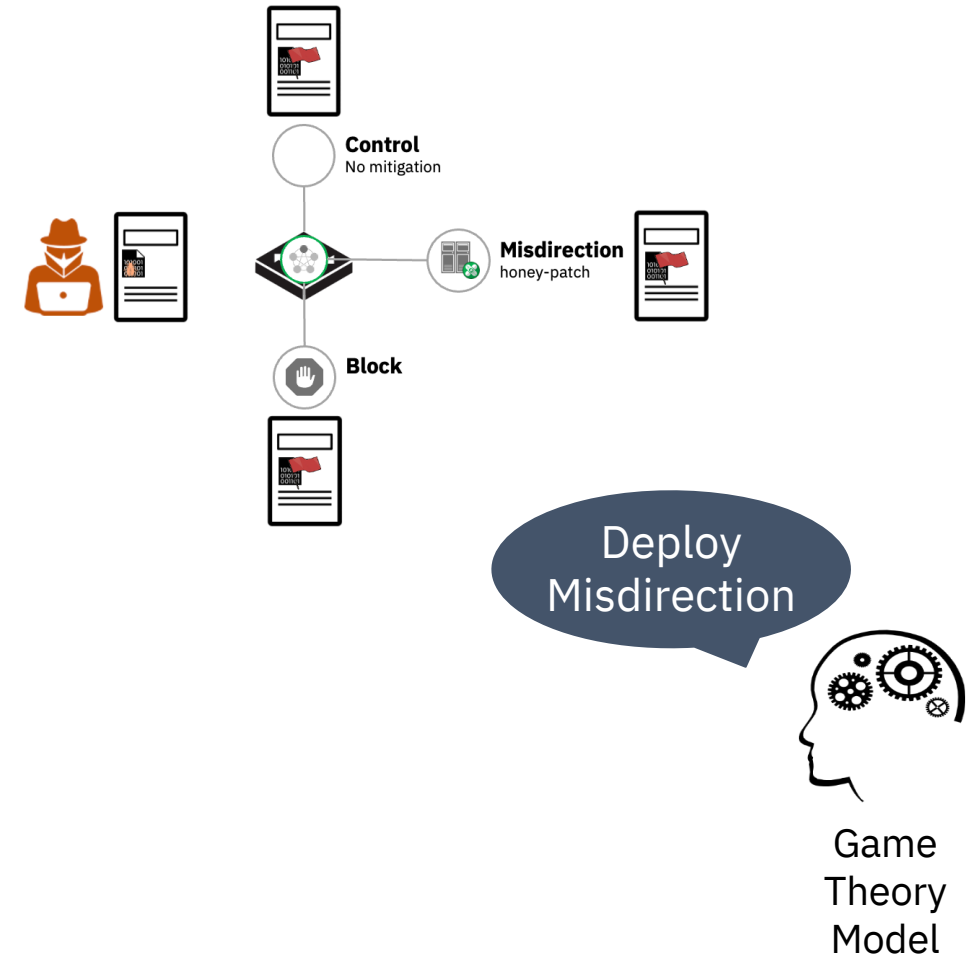
Conclusion

- Utility returns vary in the three models only for the initial 3 states of the game.
- Irrespective of parameter settings, all models recommend the defender to honey-patch the next vulnerable application.




Conclusion

- Utility returns vary in the three models only for the initial 3 states of the game.
- Irrespective of parameter settings, all models recommend the defender to honey-patch the next vulnerable application.
- Model parameters set randomly or by expert may not imitate the true model representative of the real-world attack scenario.



Summary

- ❖ Cybersecurity exercise helped gain insightful knowledge
- ❖ Closely analyzed interactions in deception-based setup
- ❖ Relaxed control over the different modalities
- ❖ Explore more ways to analyze attack behavior
- ❖ Obtain truer estimates for game-theoretic models



Thank You!